

# EN.601.414/614

# Computer Networks

## HTTP and the Web

Xin Jin

Spring 2019 (MW 3:00-4:15pm in Shaffer 301)




<https://github.com/xinjin/course-net>

# Agenda

- **HTTP and the Web**
- **Improving HTTP Performance**

# The Web: History

- **World Wide Web (WWW): a distributed database of “pages” linked through Hypertext Transport Protocol (HTTP)**
  - First HTTP implementation – 1990
    - [Tim Berners-Lee](#) at CERN
    - [Turing award](#) at 2016: for inventing the World Wide Web, the first web browser, and the fundamental protocols and algorithms allowing the Web to scale



# WWW != Internet

- **Vint Cerf, Robert Kahn**
- **Turing award at 2004:** for pioneering work on internetworking, including the design and implementation of the Internet's basic communications protocols, TCP/IP, and for inspired leadership in networking.



# The Web: History (cont'd)

- **World Wide Web (WWW): a distributed database of “pages” linked through Hypertext Transport Protocol (HTTP)**
  - HTTP/0.9 – 1991
    - Simple GET command for the Web
  - HTTP/1.0 – 1992
    - Client/server information, simple caching
  - HTTP/1.1 – 1996
    - Performance and security optimizations
  - HTTP/2 – 2015
    - Latency optimizations via request multiplexing over single TCP connection
    - Binary protocol instead of text

# Web components

- **Infrastructure:**

- Clients
- Servers (DNS, CDN, Datacenters)

- **Content:**

- URL: naming content
- HTML: formatting content

- **Protocol for exchanging information: HTTP**

# URL: Uniform Record Locator

- `protocol://host-name[:port]/directory-path/resource`

- **Extend the idea of hierarchical hostnames to include anything in a file system**

- `https://github.com/xinjin/course-net/blob/master/slides/lec01_introduction.pptx`

- **Extend to program executions as well...**

- `http://us.f413.mail.yahoo.com/ym/ShowLetter?box=%40B%40Bulk&MsgId=2604_1744106_29699_1123_1261_0_28917_3552_1289957100&Search=&Nhead=f&YY=31454&order=down&sort=date&pos=0&view=a&head=b`

- Server side processing can be included in the name

# URL: Uniform Record Locator

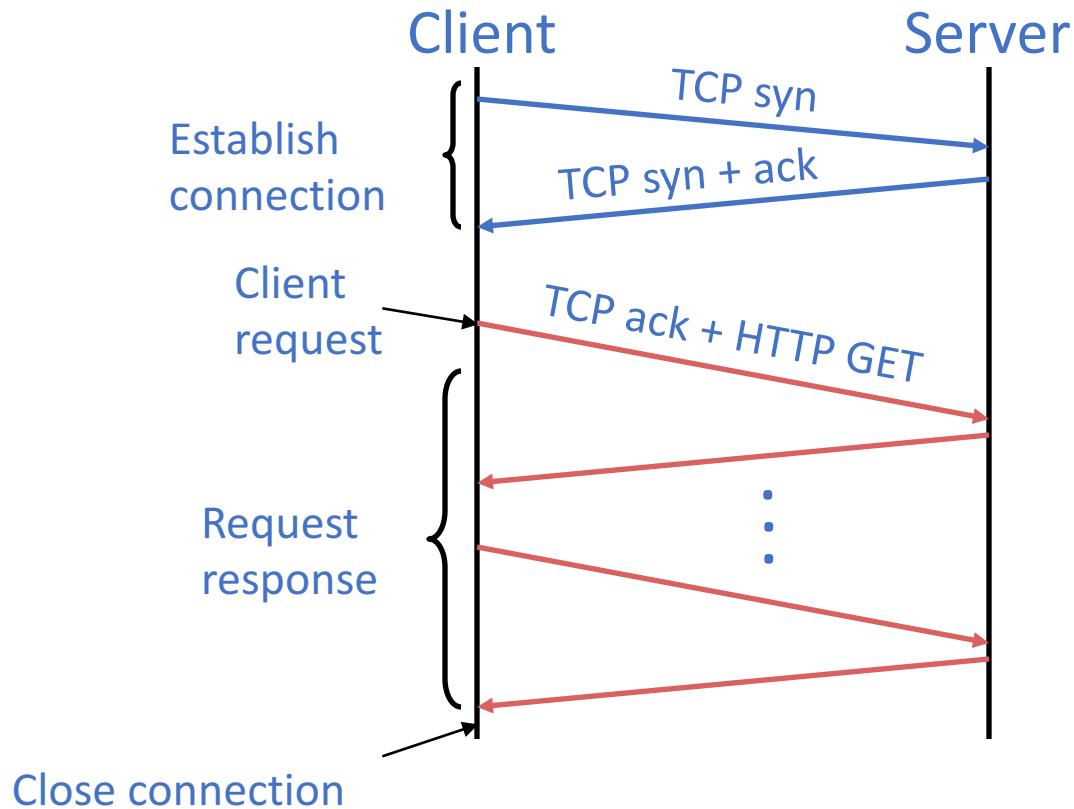
- `protocol://host-name[:port]/directory-path/resource`
  - `protocol`: http, ftp, https, smtp, rtsp, *etc.*
  - `hostname`: DNS name, IP address
  - `port`: defaults to protocol's standard port
    - *e.g.*, http: 80, https: 443
  - `directory path`: hierarchical, reflecting file system
  - `resource`: Identifies the desired resource



# Hyper Text Transfer Protocol (HTTP)

- **Client-server architecture**
  - Server is “always on” and “well known”
  - Clients initiate contact to server
- **Synchronous request/reply protocol**
  - Runs over TCP, Port 80
- **Stateless**
- **ASCII format**
  - Before HTTP/2

# Steps in HTTP request/response



# Method types (HTTP 1.1)

- **GET, HEAD**
- **POST**
  - Send information (e.g., web forms)
- **PUT**
  - Upload file in entity body to path specified in URL field
- **DELETE**
  - Delete file specified in the URL field

# Client-to-server communication

- **HTTP Request Message**

➤ Request line: method, resource, and protocol version

*request line* → **GET /somedir/page.html HTTP/1.1**

*header lines*

Host: www.someschool.edu  
User-agent: Mozilla/4.0  
Connection: close  
Accept-language: fr  
(blank line)

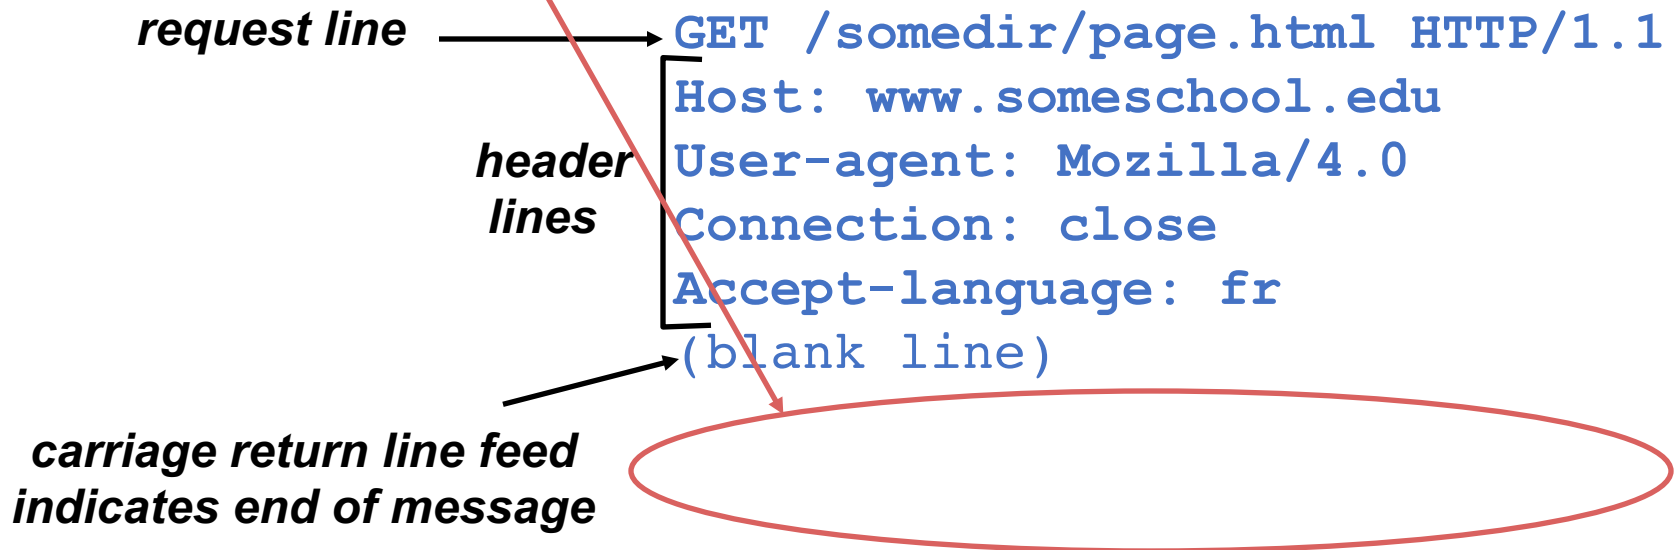
carriage return line feed indicates end of message

*carriage return line feed  
indicates end of message*

# Client-to-server communication

- **HTTP Request Message**

- Request line: method, resource, and protocol version
- Request headers: provide info or modify request
- Body: optional data (e.g., to “POST” data to server)



# Server-to-client communication

## • HTTP Response Message

- Status line: protocol version, status code, status phrase
- Response headers: provide information
- Body: optional data

### ***status line***

(protocol, status code,  
status phrase)

### ***header lines***

### ***data***

e.g., requested HTML file

HTTP/1.1 200 OK

Connection close

Date: Thu, 06 Jan 2017 12:00:15 GMT

Server: Apache/1.3.0 (Unix)

Last-Modified: Mon, 22 Jun 2006 ...

Content-Length: 6821

Content-Type: text/html

(blank line)

data data data data data ...

# HTTP is stateless

- **Each request-response treated independently**
  - Servers not required to retain state
- **Good: Improves scalability on the server-side**
  - Failure handling is easier
  - Can handle higher rate of requests
  - Order of requests doesn't matter
- **Bad: Some applications need persistent state**
  - Need to uniquely identify user or store temporary info
  - e.g., Shopping cart, user profiles, usage tracking, ...

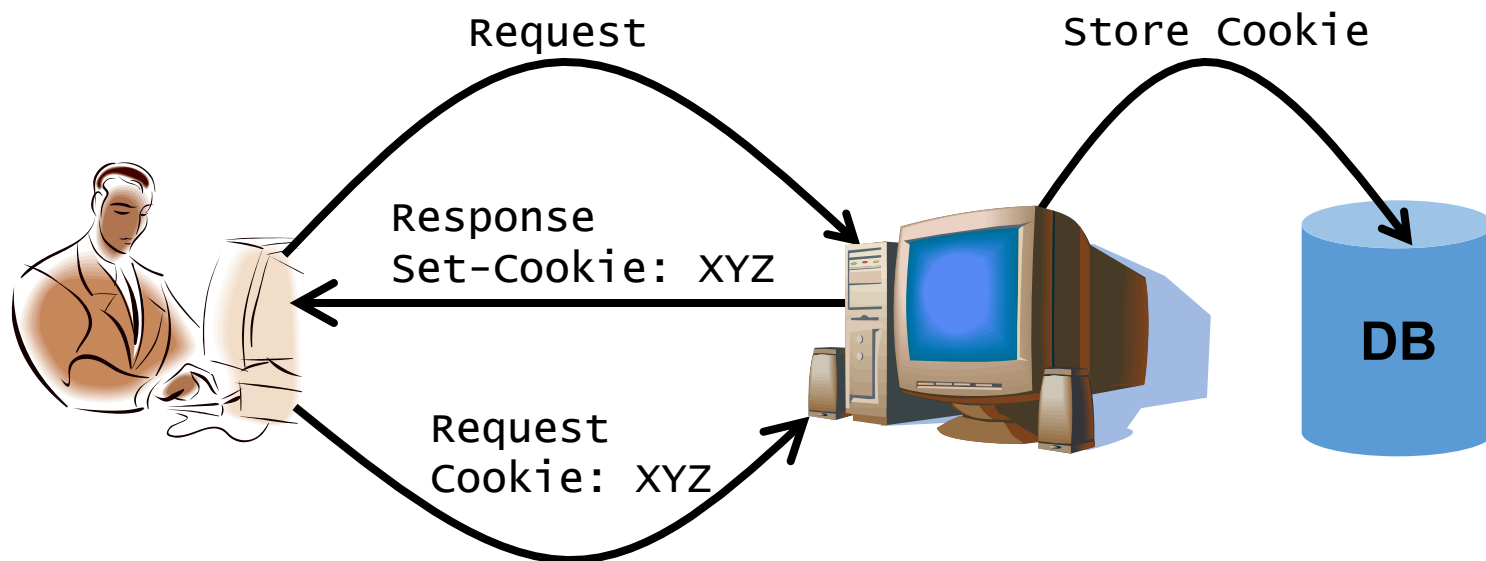
# Question

- **How does a stateless protocol keep state?**



# State in a stateless protocol: Cookies

- **Client-side state maintenance**
  - Client stores small state on behalf of server
  - Client sends state in future requests to the server
- **Can provide authentication**



# “Abuse” of cookies

- **Excellent marketing opportunities and concerns for privacy**
  - Cookies permit sites to learn a lot about you
  - You may unknowingly supply personal info to sites
  - Advertising companies tracks your preferences and viewing history across sites

# Performance goals

- **User**

- Fast downloads (not identical to low-latency communication!)
- High availability

- **Content provider**


- Happy users (hence, above)
- Cost-effective infrastructure

- **Network (secondary)**

- Avoid overload

# Solutions?

Improve networking protocols including HTTP, TCP, etc.



- **User**

- Fast downloads (not identical to low-latency communication!)
- High availability

- **Content provider**

- Happy users (hence, above)
- Cost-effective infrastructure

- **Network (secondary)**

- Avoid overload

# Solutions?

Improve networking protocols including HTTP, TCP, etc.

- **User**

- Fast downloads (not identical to low-latency communication!)
- High availability

- **Content provider**

- Happy users (hence, above)
- Cost-effective infrastructure

- **Network (secondary)**

- Avoid overload

Caching and replication

# Solutions?

- **User**

- Fast downloads (not identical to low-latency communication!)
- High availability

- **Content provider**

- Happy users (hence, above)
- Cost-effective infrastructure


- **Network (secondary)**

- Avoid overload

Improve networking protocols including HTTP, TCP, etc.

Caching and replication

Exploit economies of scale; e.g., webhosting, CDNs, datacenters



# HTTP performance

- **Most Web pages have multiple objects**
  - e.g., HTML file and a bunch of embedded images
- **How do you retrieve those objects (naively)?**
  - One item at a time
- **New TCP connection per (small) object!**


# Object request response time

- **RTT (round-trip time)**

- Time for a small packet to travel from client to server and back

- **Response time**

- 1 RTT for TCP setup
- 1 RTT for HTTP request and first few bytes
- Transmission time
- **Total** = 2RTT + Transmission Time






# Non-persistent connections

- **Default in HTTP/1.0**
- **$2RTT + \triangle$  for each object in the HTML file!**
  - One more  $2RTT + \triangle$  for the HTML file itself
- **Doing the same thing over and over again**
  - Inefficient

# Concurrent requests and responses

- Use multiple connections in parallel
- Does not necessarily maintain order of responses
- Client = 😊
- Content provider = 😊
- Network = 😞 Why?

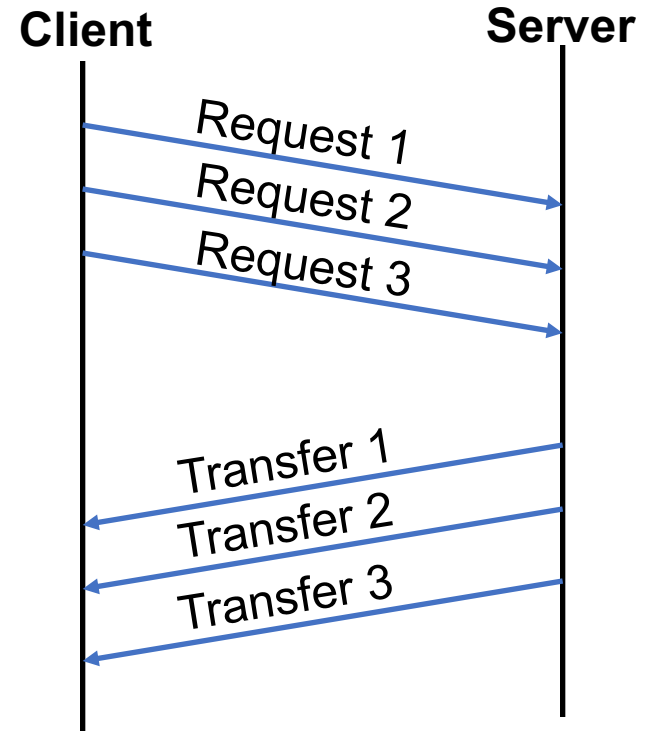


# Persistent connections

- **Maintain TCP connection across multiple requests**
  - Including transfers subsequent to current page
  - Client or server can tear down connection
- **Advantages**
  - Avoid overhead of connection set-up and tear-down
  - Allow underlying layers (e.g., TCP) to learn about RTT and bandwidth characteristics
- **Default in HTTP/1.1**

# Pipelined requests & responses

- Batch requests and responses to reduce the number of packets
- Multiple requests can be contained in one TCP segment



# Scorecard: Getting $n$ small objects

- Time dominated by latency
- One-at-a-time:  $\sim 2n$  RTT
- $m$  concurrent:  $\sim 2\lceil n/m \rceil$  RTT
- Persistent:  $\sim (n+1)$ RTT
- Pipelined:  $\sim 2$  RTT
- Pipelined/Persistent:  $\sim 2$  RTT first time, RTT later

# Scorecard: Getting $n$ large objects each of size $F$

- Time dominated by bandwidth
- **One-at-a-time:**  $\sim nF/B$
- **$m$  concurrent:**  $\sim [n/m] F/B$ 
  - Assuming shared with large population of users and each TCP connection gets the same bandwidth
- **Pipelined and/or persistent:**  $\sim nF/B$ 
  - The only thing that helps is getting more bandwidth

# Summary

- **HTTP/1.1**
  - Text-based protocol
  - Being replaced by binary HTTP/2 protocol
- **Many ways to improve performance**
  - Pipelining and batching
- **Assignment 1 is due next Friday**

Thanks!  
Q&A