

EN.601.414/614

Computer Networks

HTTP and the Web

Xin Jin

Fall 2020 (TuTh 1:30-2:45pm on Zoom)




<https://github.com/xinjin/course-net>

Agenda

- **HTTP and the Web**
- **Improving HTTP Performance**

The Web: History

- **World Wide Web (WWW): a distributed database of “pages” linked through Hypertext Transport Protocol (HTTP)**
 - First HTTP implementation – 1990
 - [Tim Berners-Lee](#) at CERN
 - [Turing award](#) at 2016: for inventing the World Wide Web, the first web browser, and the fundamental protocols and algorithms allowing the Web to scale



WWW != Internet

- **Vint Cerf, Robert Kahn**
- **Turing award at 2004:** for pioneering work on internetworking, including the design and implementation of the Internet's basic communications protocols, TCP/IP, and for inspired leadership in networking.



The Web: History (cont'd)

- **World Wide Web (WWW): a distributed database of “pages” linked through Hypertext Transport Protocol (HTTP)**
 - HTTP/0.9 – 1991
 - Simple GET command for the Web
 - HTTP/1.0 – 1992
 - Client/server information, simple caching
 - HTTP/1.1 – 1996
 - Performance and security optimizations
 - HTTP/2 – 2015
 - Latency optimizations via request multiplexing over single TCP connection
 - Binary protocol instead of text

Web components

- **Infrastructure:**

- Clients
- Servers (DNS, CDN, Datacenters)

- **Content:**

- URL: naming content
- HTML: formatting content

- **Protocol for exchanging information: HTTP**

URL: Uniform Record Locator

- `protocol://host-name[:port]/directory-path/resource`

- **Extend the idea of hierarchical hostnames to include anything in a file system**

- `https://github.com/xinjin/course-net/blob/master/slides/lec01_introduction.pptx`

- **Extend to program executions as well...**

- `http://us.f413.mail.yahoo.com/ym/ShowLetter?box=%40B%40Bulk&MsgId=2604_1744106_29699_1123_1261_0_28917_3552_1289957100&Search=&Nhead=f&YY=31454&order=down&sort=date&pos=0&view=a&head=b`

- Server side processing can be included in the name

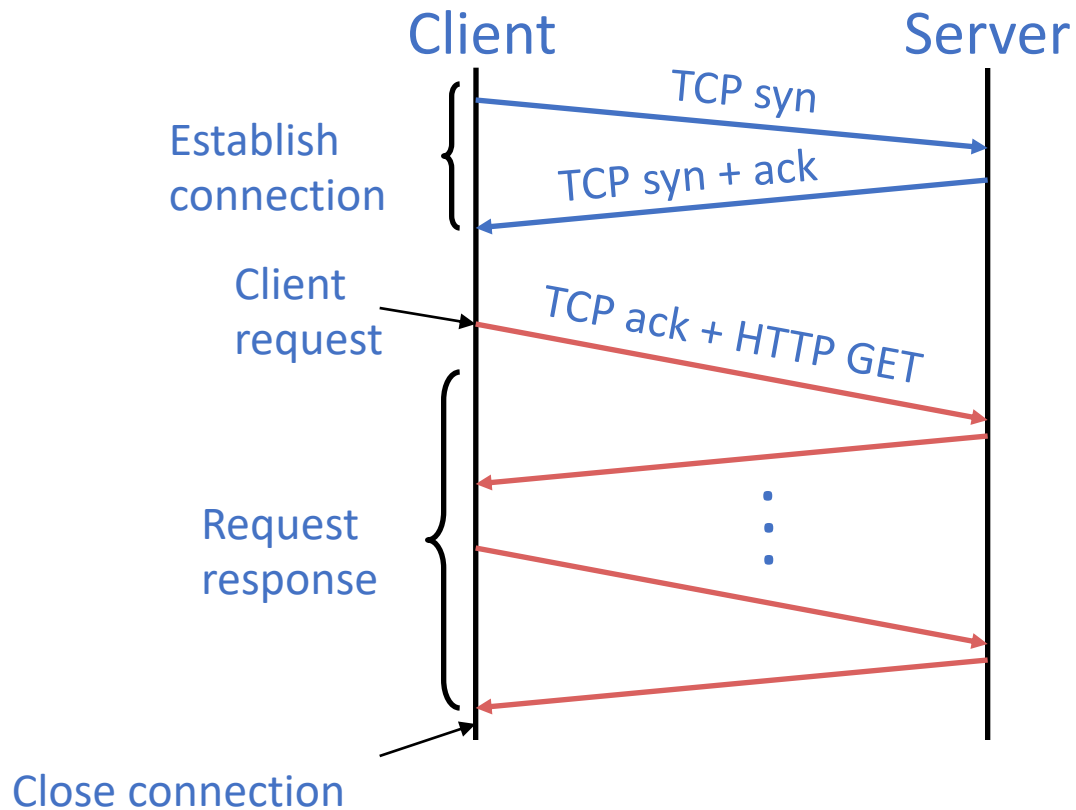
URL: Uniform Record Locator

- `protocol://host-name[:port]/directory-path/resource`
 - `protocol`: http, ftp, https, smtp, rtsp, *etc.*
 - `hostname`: DNS name, IP address
 - `port`: defaults to protocol's standard port
 - *e.g.*, http: 80, https: 443
 - `directory path`: hierarchical, reflecting file system
 - `resource`: Identifies the desired resource

Hyper Text Transfer Protocol (HTTP)

- **Client-server architecture**
 - Server is “always on” and “well known”
 - Clients initiate contact to server
- **Synchronous request/reply protocol**
 - Runs over TCP, Port 80
- **Stateless**
- **ASCII format**
 - Before HTTP/2

Steps in HTTP request/response



Method types (HTTP 1.1)

- **GET, HEAD**
- **POST**
 - Send information (e.g., web forms)
- **PUT**
 - Upload file in entity body to path specified in URL field
- **DELETE**
 - Delete file specified in the URL field

Client-to-server communication

- **HTTP Request Message**

➤ Request line: method, resource, and protocol version

request line → **GET /somedir/page.html HTTP/1.1**

header lines →
Host: www.someschool.edu
User-agent: Mozilla/4.0
Connection: close
Accept-language: fr
(blank line)

*carriage return (\r) line feed (\n)
indicates end of message*

Client-to-server communication

- **HTTP Request Message**

- Request line: method, resource, and protocol version
- Request headers: provide information
- Body: optional data (e.g., to “POST” data to server)

request line → GET /somedir/page.html HTTP/1.1

header lines [Host: www.someschool.edu
User-agent: Mozilla/4.0
Connection: close
Accept-language: fr
(blank line)

→ (blank line)

*carriage return (\r) line feed (\n)
indicates end of message*

The diagram illustrates the structure of an HTTP request message. It consists of three main parts: the request line, header lines, and a blank line indicating the end of the message. The request line is 'GET /somedir/page.html HTTP/1.1'. The header lines are 'Host: www.someschool.edu', 'User-agent: Mozilla/4.0', 'Connection: close', and 'Accept-language: fr'. A red circle highlights the body area, which is currently empty. A red arrow points from the text 'carriage return (\r) line feed (\n) indicates end of message' to the blank line. Another red arrow points from the text 'Body: optional data (e.g., to “POST” data to server)' to the body area.

Server-to-client communication

• HTTP Response Message

- Status line: protocol version, status code, status phrase
- Response headers: provide information
- Body: optional data

status line

(protocol, status code, status phrase)

header lines

data

e.g., requested HTML file

HTTP/1.1 200 OK

Connection close

Date: Thu, 06 Jan 2017 12:00:15 GMT

Server: Apache/1.3.0 (Unix)

Last-Modified: Mon, 22 Jun 2006 ...

Content-Length: 6821

Content-Type: text/html

(blank line)

data data data data data ...

HTTP is stateless

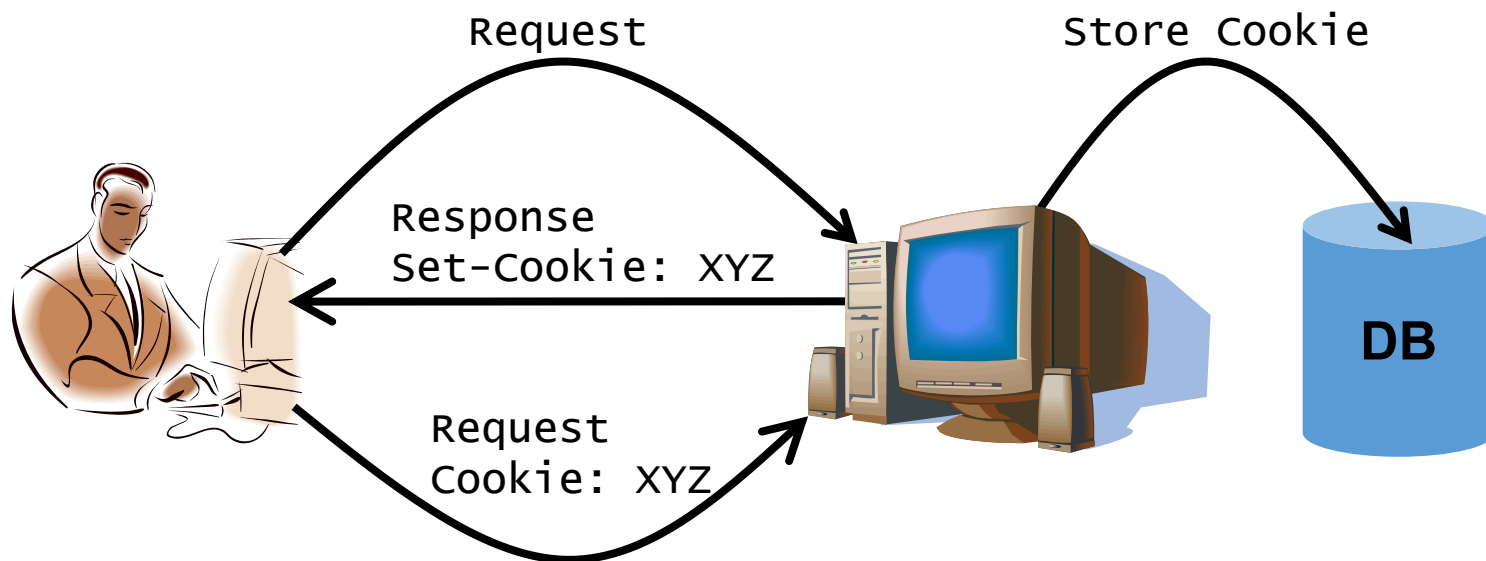
- **Each request-response treated independently**
 - Servers not required to retain state
- **Good: Improves scalability on the server-side**
 - Failure handling is easier
 - Can handle higher rate of requests
 - Order of requests doesn't matter
- **Bad: Some applications need persistent state**
 - Need to uniquely identify user or store temporary info
 - e.g., shopping cart, user profiles, usage tracking, ...

Question

- **How does a stateless protocol keep state?**

State in a stateless protocol: Cookies

- **Client-side state maintenance**
 - Client stores small state on behalf of server
 - Client sends state in future requests to the server
- **Can provide authentication**



“Abuse” of cookies

- **Excellent marketing opportunities and concerns for privacy**
 - Cookies permit sites to learn a lot about you
 - You may unknowingly supply personal info to sites
 - Advertising companies tracks your preferences and viewing history across sites

Performance goals

- **User**

- Fast downloads (not identical to low-latency communication!)
- High availability

- **Content provider**


- Happy users (hence, above)
- Cost-effective infrastructure

- **Network (secondary)**

- Avoid overload

Solutions?

Improve networking protocols including HTTP, TCP, etc.



- **User**

- Fast downloads (not identical to low-latency communication!)
- High availability

- **Content provider**

- Happy users (hence, above)
- Cost-effective infrastructure

- **Network (secondary)**

- Avoid overload

Solutions?

Improve networking protocols including HTTP, TCP, etc.

- **User**

- Fast downloads (not identical to low-latency communication!)
- High availability

- **Content provider**

- Happy users (hence, above)
- Cost-effective infrastructure

- **Network (secondary)**

- Avoid overload

Caching and replication

Solutions?

- **User**

- Fast downloads (not identical to low-latency communication!)
- High availability

- **Content provider**

- Happy users (hence, above)
- Cost-effective infrastructure

- **Network (secondary)**

- Avoid overload

Improve networking protocols including HTTP, TCP, etc.

Caching and replication

Exploit economies of scale; e.g., webhosting, CDNs, datacenters



HTTP performance

- **Most Web pages have multiple objects**
 - e.g., HTML file and a bunch of embedded images
- **How do you retrieve those objects (naively)?**
 - One item at a time
- **New TCP connection per (small) object!**


Object request response time

- **RTT (round-trip time)**

- Time for a small packet to travel from client to server and back


- **Response time**

- 1 RTT for TCP setup
- 1 RTT for HTTP request and first few bytes
- Transmission time
- **Total** = 2RTT + Transmission Time



Concurrent requests and responses

- Use multiple connections in parallel
- Does not necessarily maintain order of responses
- Client = 😊
- Content provider = 😊
- Network = 😞 Why?



Non-persistent connections

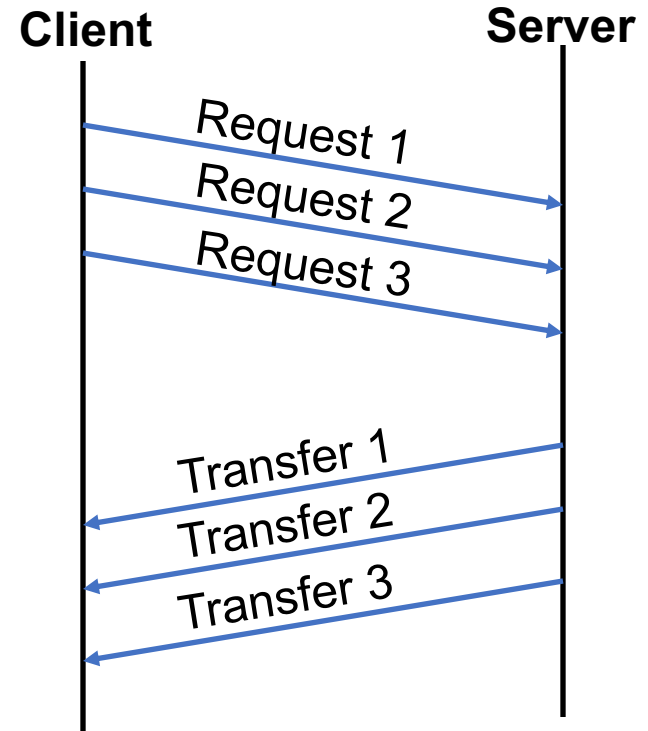
- **Default in HTTP/1.0**
- **2RTT+ Δ for each object in the HTML file!**
 - One more 2RTT+ Δ for the HTML file itself
- **Doing the same thing over and over again**
 - Inefficient

Persistent connections

- **Maintain TCP connection across multiple requests**
 - Including transfers subsequent to current page
 - Client or server can tear down connection
- **Advantages**
 - Avoid overhead of connection set-up and tear-down
 - Allow underlying layers (e.g., TCP) to learn about RTT and bandwidth characteristics
- **Default in HTTP/1.1**

Pipelined requests & responses

- Batch requests and responses to reduce the number of packets
- Multiple requests can be contained in one TCP segment



Scorecard: Getting n small objects

- Time dominated by latency
- One-at-a-time: $\sim 2n$ RTT
- m concurrent: $\sim 2\lceil n/m \rceil$ RTT
- Persistent: $\sim (n+1)$ RTT
- Pipelined: ~ 2 RTT
- Pipelined/Persistent: ~ 2 RTT first time, RTT later

Scorecard: Getting n large objects each of size F

- Time dominated by bandwidth
- **One-at-a-time:** $\sim nF/B$
- **m concurrent:** $\sim [n/m] F/B$
 - Assuming shared with large population of users and each TCP connection gets the same bandwidth
- **Pipelined and/or persistent:** $\sim nF/B$
 - The only thing that helps is getting more bandwidth

Group Discussion

- **Topic: web performance optimization**

- What are the common techniques for optimizing HTTP performance?
- Pick an Internet application. Which techniques are useful in improving its performance? Which are not?

- **Discuss in groups, and each group chooses a leader to summarize the discussion**

- In your group discussion, please do not dominate the discussion, and give everyone a chance to speak

Summary

- **HTTP/1.1**
 - Text-based protocol
 - Being replaced by binary HTTP/2 protocol
- **Many ways to improve performance**
 - Concurrent connections, persistent connections, pipelining
- **Assignment 1 is due next Sunday**

Thanks!
Q&A