

EN.601.414/614

Computer Networks

CDN and DNS

Xin Jin

Spring 2019 (MW 3:00-4:15pm in Shaffer 301)



<https://github.com/xinjin/course-net>

Question on differences between connection and circuit

- **Connection**

- A transport layer concept
- Resources are reserved at end hosts (sender & receiver)
- Need the underlying network layer to send data

- **Circuit**

- A network layer concept
- Resources are reserved at each hop
- Circuit switching is one way to support a connection; packet switching is the other way

Agenda

- **CDN: Content Distribution Network**
- **DNS: Domain Name System**

Recap:

Improving HTTP performance

- **Optimizing connections using three “P”s**
 - Persistent connections
 - Parallel/concurrent connections
 - Pipelined transfers over the same connection

Caching

- **Why does caching work?**

- Exploit locality of reference

- **How well does caching work?**

- Very well, up to a limit

- Large overlap in content

- But many unique requests

- Empirical result: effectiveness of caching (cache hit ratio) grows logarithmically with user size

Caching: How

- **Modifier to GET requests:**

- **If-modified-since** – returns “not modified” if resource not modified since specified time

```
GET /somedir/page.html HTTP/1.1
```

```
Host: www.someschool.edu
```

```
User-agent: Mozilla/4.0
```

```
If-modified-since: Wed, 18 Jan 2017 10:25:50 GMT  
(blank line)
```

Caching: How

- **Modifier to GET requests:**
 - **If-modified-since** – returns “not modified” if resource not modified since specified time
- **Client specifies “if-modified-since” time in request**
- **Server compares this against “last modified” time of resource**
- **Server returns “Not Modified” if resource has not changed**
- **.... or a “OK” with the latest version otherwise**

Caching: How

- **Modifier to GET requests:**

- **If-modified-since** – returns “not modified” if resource not modified since specified time

- **Response header:**

- **Expires** – how long it's safe to cache the resource

- **No-cache** – ignore all caches; always get resource directly from server


Caching: Where?

- **Options**

- Client (browser)
- Forward proxies
- Reverse proxies
- Content Distribution Network

Caching: Where?

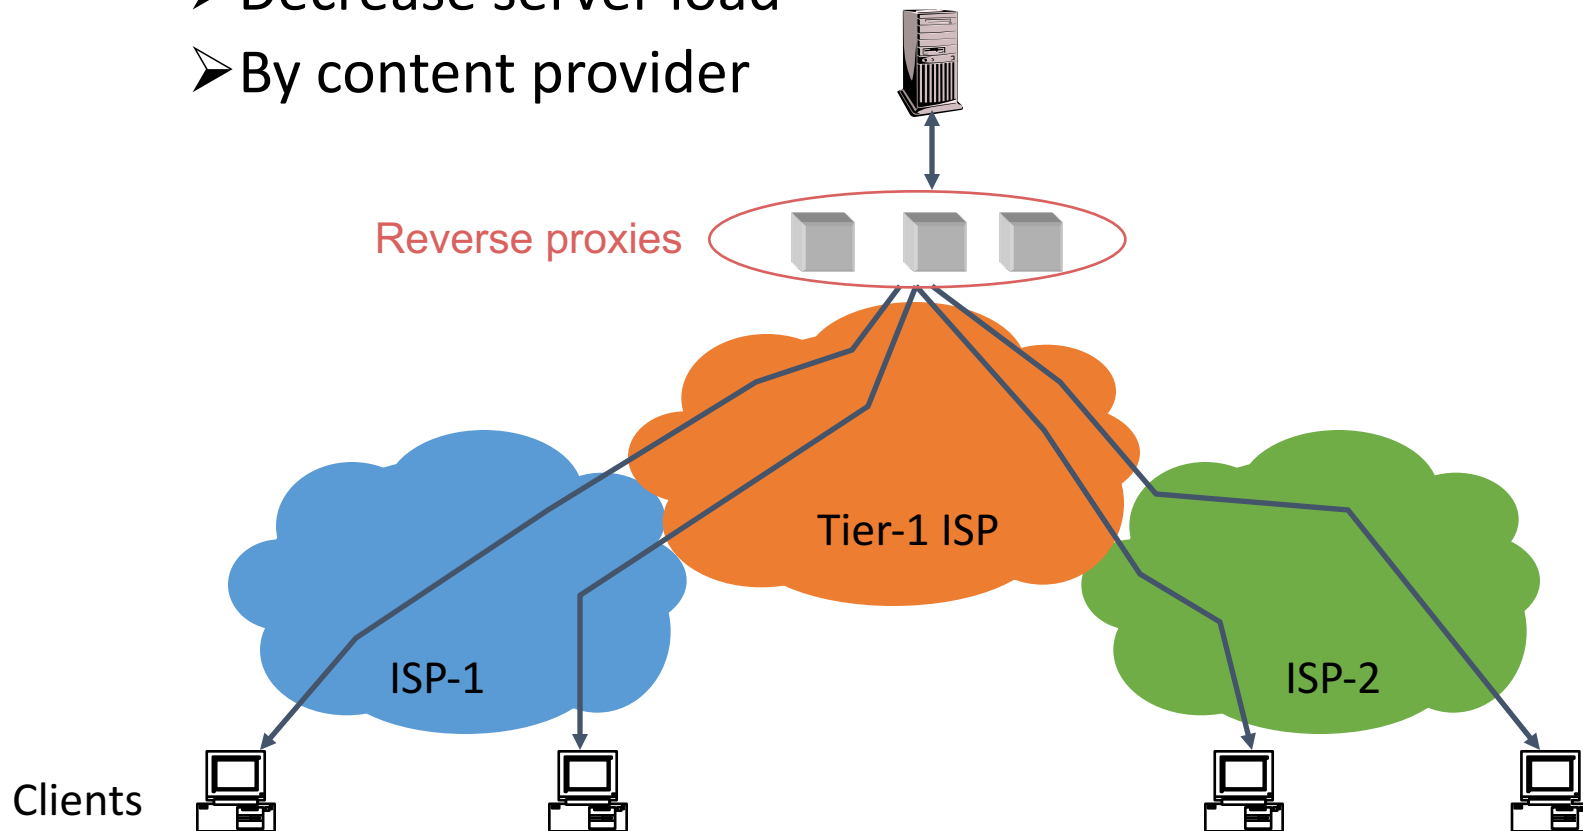
- **Many clients transfer same information**
 - Generate unnecessary server and network load
 - Clients experience unnecessary latency



Caching with Reverse Proxies


- **Cache documents close to server**

- Decrease server load
- By content provider



Caching with Forward Proxies

- **Cache documents close to clients**
 - Reduce network traffic and decrease latency
 - By ISPs or enterprises



Replication

- **Replicate popular Websites across many machines**
 - Spread load across servers
 - Place content closer to clients
 - Help when content isn't cacheable

Recap:

Improving HTTP performance

- **Optimizing connections using three “P”s**
 - Persistent connections
 - Parallel/concurrent connections
 - Pipelined transfers over the same connection
- **Caching**
 - Forward proxy: close to clients
 - Reverse proxy: close to servers
- **Replication**

Content Distribution Networks (CDN)

- **Caching and replication as a service**
- **Large-scale distributed storage infrastructure (usually) administered by one entity**
 - e.g., Akamai has servers in 20,000+ locations
- **Combination of caching and replication**
 - **Pull**: Direct result of clients' requests (caching)
 - **Push**: Expectation of high access rate (replication)
- **Can do some processing to handle dynamic webpage content**

Cost-effective content delivery

- **General theme: multiple sites hosted on shared physical infrastructure**
 - Efficiency of statistical multiplexing
 - Economies of scale (volume pricing, etc.)
 - Amortization of human operator costs
- **Examples:**
 - CDNs
 - Web hosting companies
 - Cloud infrastructure

CDN example – Akamai

- **Akamai creates new domain names for each client**
 - e.g., a128.g.akamai.net for cnn.com
- **The client content provider modifies content so that embedded URLs reference new domains**
 - “Akamaize” content
 - e.g., <http://www.cnn.com/image-of-the-day.gif> becomes <http://a128.g.akamai.net/image-of-the-day.gif>
- **Requests now sent to CDN’s infrastructure**

How to direct clients to particular replicas?

- **In order to**

- Balancing load across server replicas
- Pairing clients with nearby servers to decrease latency and overall bandwidth usage

DNS: Domain name system

Internet names & addresses

- **Machine addresses: e.g., 141.212.113.143**
 - Router-usable labels for machines
 - Conforms to network structure (the “where”)
- **Machine names: e.g., cs.jhu.edu**
 - Human-usable labels for machines
 - Conforms to organizational structure (the “who”)
- **The Domain Name System (DNS) is how we map from one to the other**
 - A directory service

Why?

- **Convenience**

- Easier to remember www.google.com than 216.58.216.100

- **Provides a level of indirection!**

- Decoupled names from addresses
 - Many uses beyond just naming a specific host

DNS: History

- **Initially all host-address mappings were in a `hosts.txt` file (in `/etc/hosts`):**
 - Maintained by the Stanford Research Institute (SRI)
 - Changes were submitted to SRI by email
 - New versions of `hosts.txt` periodically FTP'd from SRI

DNS: History (cont'd)

- **As the Internet grew this system broke down**
 - SRI couldn't handle the load
 - Names were not unique
 - Hosts had inaccurate copies of `hosts.txt`
- **The Domain Name System (DNS) was invented to fix this**

Goals

- **Uniqueness: no naming conflicts**
- **Scalable**
 - Many names and frequent updates
- **Distributed, autonomous administration**
 - Ability to update my own (machines') names
 - Don't have to track everybody's updates
- **Highly available**
- **Lookups are fast**
- **Perfect consistency is a non-goal**

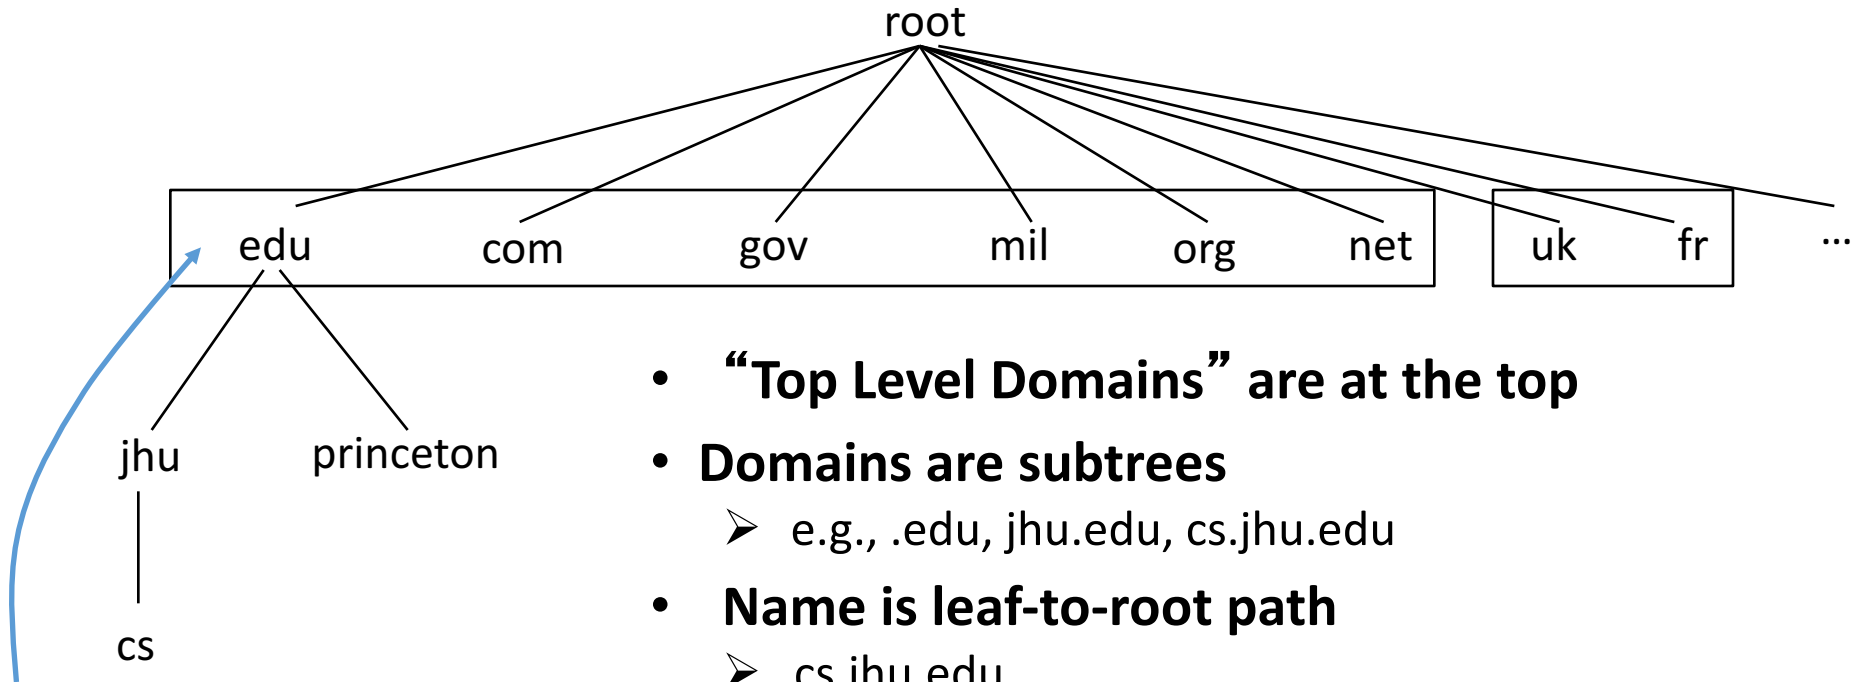
How?

- **Partition the namespace**
- **Distribute administration of each partition**
 - Autonomy to update my own (machines') names
 - Don't have to track everybody's updates
- **Distribute name resolution for each partition**
- **How should we partition things?**

Key idea: Hierarchy

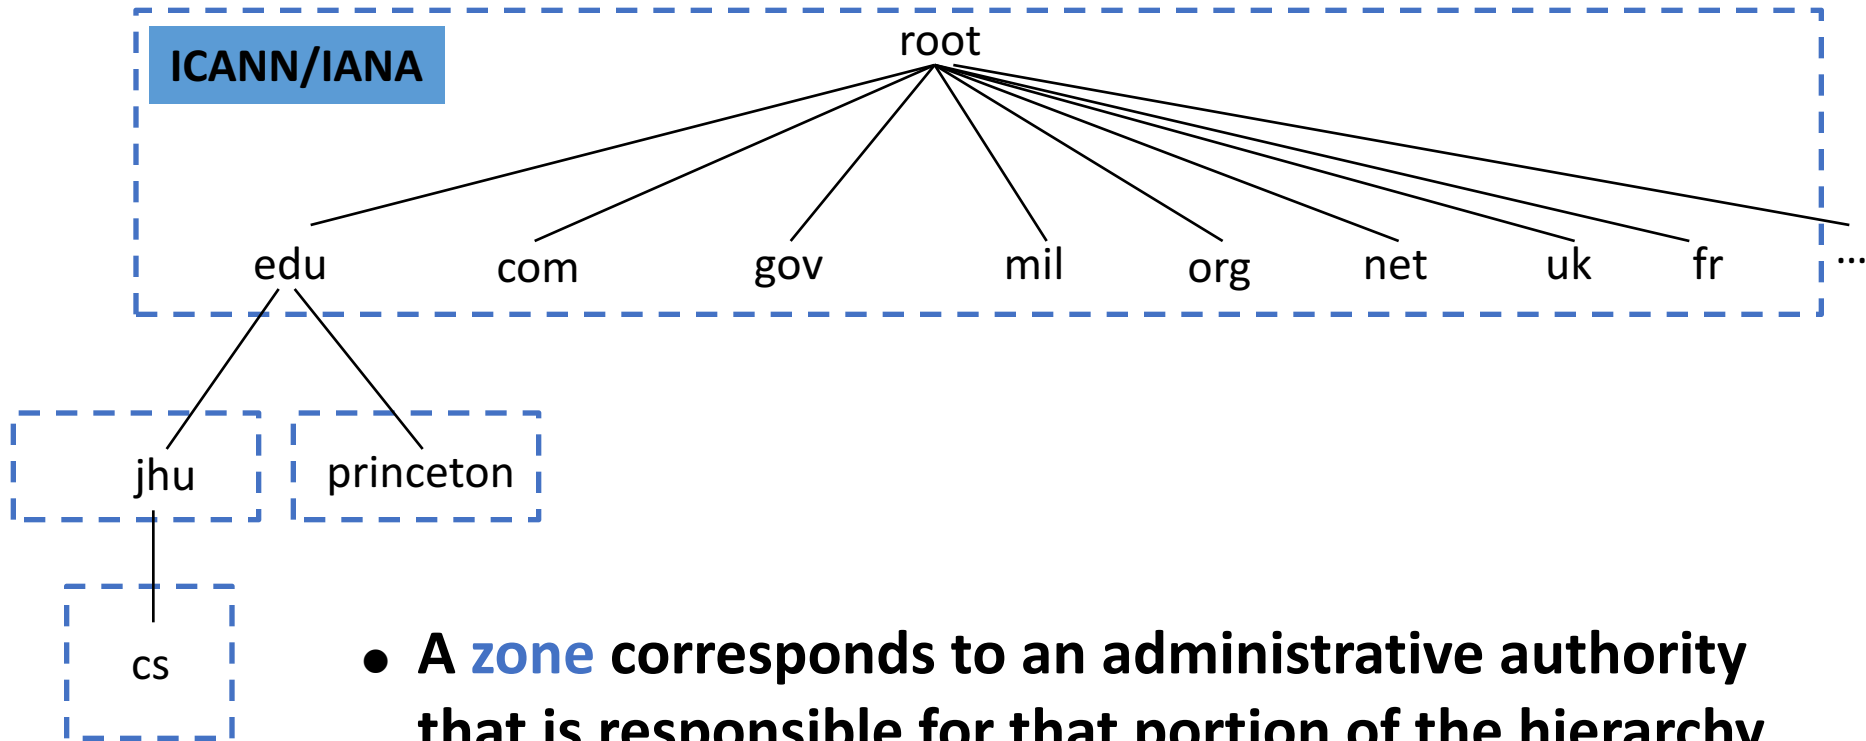
- **Three intertwined hierarchies**
 - Hierarchical namespace
 - As opposed to original flat namespace
 - Hierarchically administered
 - As opposed to centralized
 - (Distributed) hierarchy of servers
 - As opposed to centralized storage

Hierarchical namespace



- **“Top Level Domains” are at the top**
- **Domains are subtrees**
 - e.g., .edu, jhu.edu, cs.jhu.edu
- **Name is leaf-to-root path**
 - cs.jhu.edu
- **Depth of tree is arbitrary (limit 128)**
- **Name collisions trivially avoided**
 - Each domain is responsible

Hierarchical administration



Server hierarchy

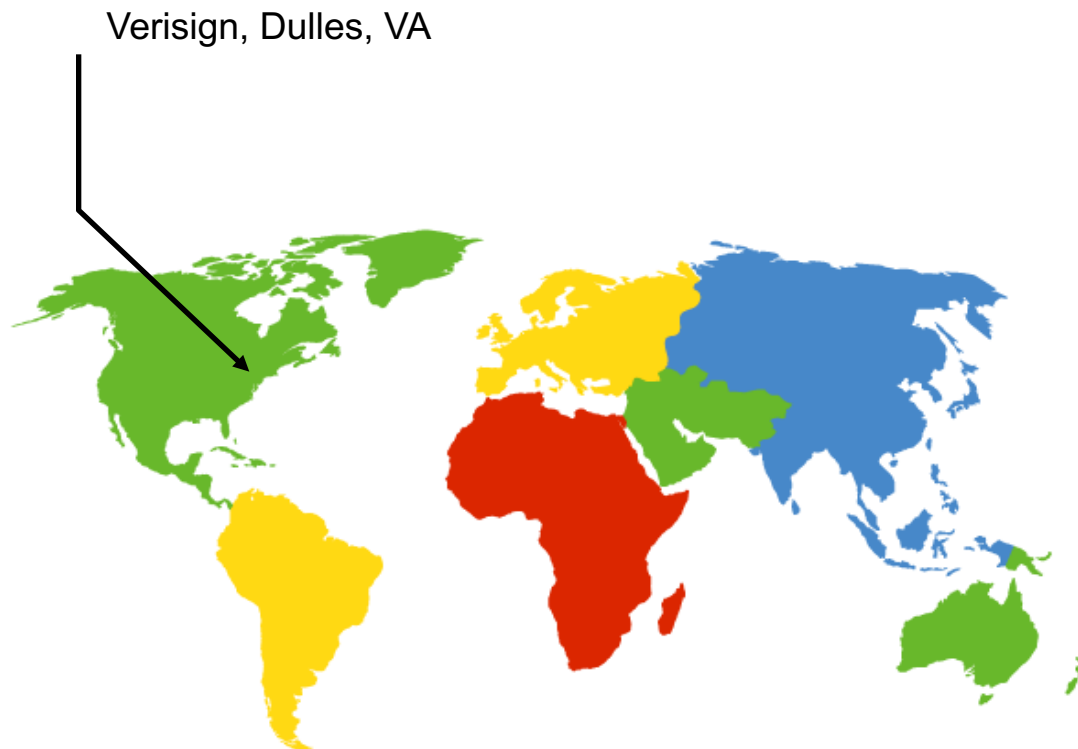
- **Top of hierarchy: Root servers**
 - Location hardwired into other servers
- **Next Level: Top-level domain (TLD) servers**
 - .com, .edu, etc.
 - Managed professionally
- **Bottom Level: Authoritative DNS servers**
 - Actually store the name-to-address mapping
 - Maintained by the corresponding administrative authority

Server hierarchy

- Each server stores a (small!) subset of the total DNS database
- An authoritative DNS server stores “**resource records**” for all DNS names in the domain that it has authority for
- Each server needs to know other servers that are responsible for the other portions of the hierarchy
 - Every server knows the root
 - Root server knows about all top-level domains

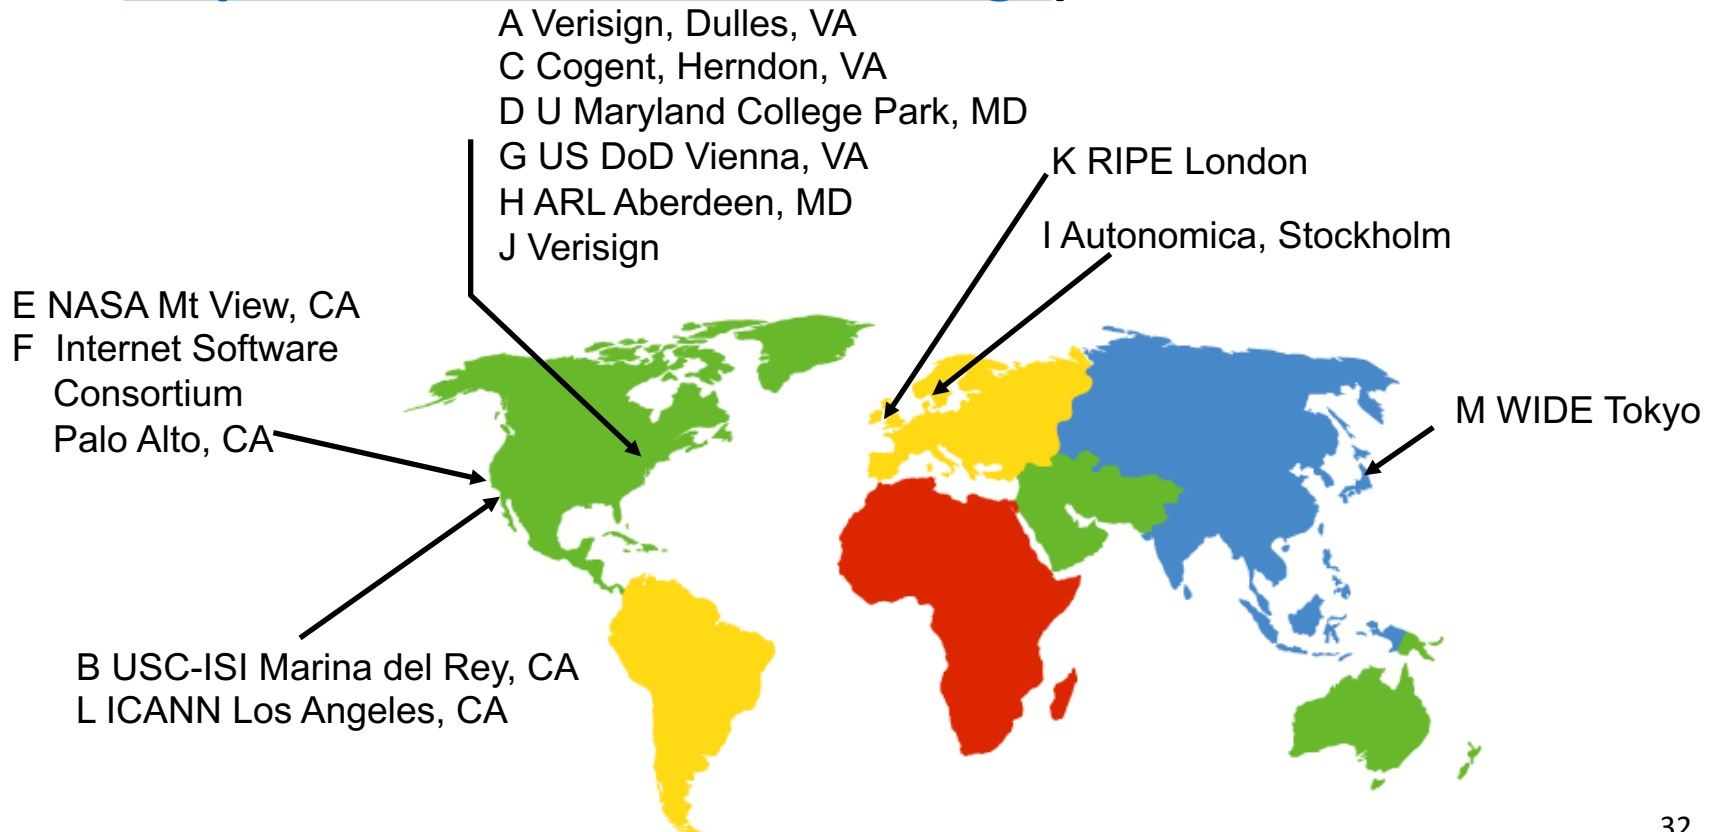
DNS root

- Located in Virginia, USA
- How do we make the root scale?



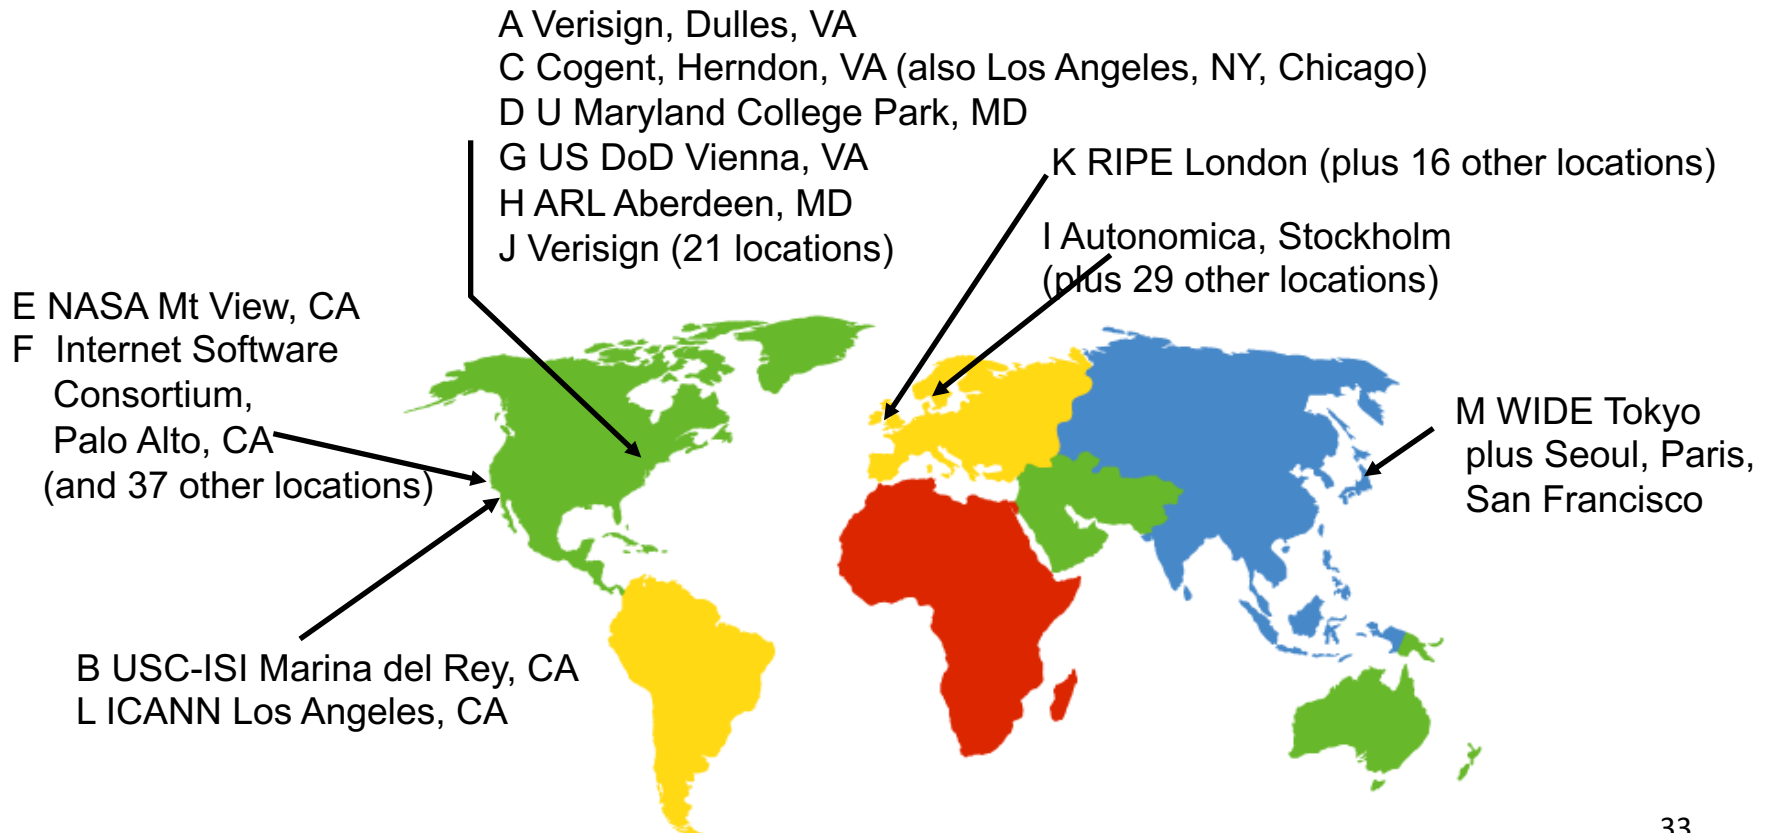
DNS root servers

- **13 root servers (labeled A-M; see <http://www.root-servers.org/>)**



DNS root servers

- **13 root servers replicated via anycast**



Anycast in a nutshell

- **Routing finds shortest paths to destination**
- **If several locations are given the same address, then the network will deliver the packet to the closest location with that address**
- **Characteristics**
 - Very robust
 - Requires no modification to routing algorithms

DNS records

- **DNS servers store resource records (RRs)**
 - RR is (name, value, type, TTL)
- **Type = A: (→ Address)**
 - name = hostname
 - value = IP address
- **Type = NS: (→ Name Server)**
 - name = domain
 - value = name of DNS server for domain

DNS records (cont'd)

- **Type = CNAME: (→ Canonical Name)**

- name = alias name for some “canonical” (real) name
 - e.g., documents.example.com is really docs.example.com
- value = canonical name

- **Type = MX: (→ Mail eXchanger)**

- name = domain in email address
- value = name(s) of mail server(s)

Inserting Resource Records into DNS

- **Register foobar.com at registrar (GoDaddy)**
 - Provide registrar with names and IP addresses of your authoritative name server(s)
 - Registrar inserts RR pairs into the .com TLD server:
 - (foobar.com, dns1.foobar.com, NS)
 - (dns1.foobar.com, 212.44.9.129, A)
- **Store resource records in your server
dns1.foobar.com**
 - e.g., type A record for www.foobar.com
 - e.g., type MX record for foobar.com

Using DNS (Client/App View)

- **Two components**

- Local DNS servers
- Resolver software on hosts


- **Local DNS server (“default name server”)**

- Clients configured with default server’s address OR learn it via a host configuration protocol (e.g., DHCP)


- **Client application**

- Obtain DNS name (e.g., from URL)
- Do `gethostbyname()` to trigger DNS request to its local DNS server


Name resolution




Name resolution




Name resolution




Name resolution



Name resolution: Recursive



Name resolution: Iterative



Two ways to resolve a name

- **Recursive name resolution**
 - Ask server to do it for you
- **Iterative name resolution**
 - Ask server who to ask next
- **The iterative example we saw is a mix of both!**

DNS protocol

- **Query and Reply messages; both with the same message format**
 - Header: identifier, flags, etc.
 - Plus resource records
 - See text for details
- **Client–server interaction on UDP Port 53**
 - Spec supports TCP too, but not always implemented

Goals: Are we there yet?

- **Uniqueness: No naming conflicts**
- **Scalable**
- **Distributed, autonomous administration**
- **Highly available?**

Reliability

- **Replicated DNS servers (primary/secondary)**
 - Name service available if at least one replica is up
 - Queries can be load-balanced between replicas
- **Usually, UDP used for queries**
 - Reliability, if needed, must be implemented on UDP
- **Try alternate servers on timeout**
 - Exponential backoff when retrying same server
- **Same identifier for all queries**
 - Don't care which server responds

Goals: Are we there yet?

- **Uniqueness: No naming conflicts**
- **Scalable**
- **Distributed, autonomous administration**
- **Highly available**
- **Fast lookups?**

DNS caching

- **Performing all these queries takes time**
 - Up to 1-second latency before starting download
- **Caching can greatly reduce overhead**
 - The top-level servers very rarely change
 - Popular sites (e.g., www.cnn.com) visited often
 - Local DNS server often has the information cached
- **How DNS caching works**
 - DNS servers cache responses to queries
 - Responses include a “time to live” (TTL) field
 - Server deletes cached entry after TTL expires

Negative caching

- **Remember things that don't work**

- Misspellings like `www.cnn.comm` and `www.cnnn.com`
- These can take a long time to fail the first time
- Good to remember that they don't work so the failure takes less time the next time around

- **Negative caching is optional**

- Not widely implemented

Important properties of DNS

- **Administrative delegation and hierarchy enables:**
 - Easy unique naming
 - “Fate sharing” for network failures
 - Reasonable trust model
 - Caching increases scalability and performance

DNS provides indirection

- **Addresses can change underneath**
 - Move `www.cnn.com` to `4.125.91.21`
- **Name could map to multiple IP addresses**
 - Load-balancing ([CDN](#))
 - Reducing latency by picking nearby servers ([CDN](#))
- **Multiple names for the same address**
 - E.g., many services (mail, www) on same machine
 - E.g., aliases like `www.cnn.com` and `cnn.com`
- **This flexibility applies only within domain!**

Summary

- **CDNs improve web performance**
 - Via replication and caching
 - Good server selection
- **DNS allows us to go to webpages without having to memorize IP addresses**
 - Allows a level of indirection that enables many functionalities including CDN server selection
- **Exercise and lab session this Wednesday**
- **Assignment 1 is due this Friday**

Thanks!
Q&A