

EN.601.414/614

Computer Networks

Routing Fundamentals

Xin Jin

Spring 2019 (MW 3:00-4:15pm in Shaffer 301)



Research Opportunities

- **You are very welcome!**

- Undergraduate: recommendation letters for applying to graduate programs
(Berkeley/Princeton/MIT/Stanford/CMU/Harvard, etc.)
- Master: join my research group as PhDs
- PhD: GBO projects

- **Requirements:**

- Get A/A+ from 601.414/614 Computer Networks
- Good programming skills in C/C++/Python
- Having taken 601.714 Advanced Computer Networks
(offered every fall) is a plus

Research Opportunities

- **Projects**

- Hardware-software co-design: leverage programmable hardware (e.g., programmable switches, GPUs, FPGAs) to improve computation, networking and storage for cloud computing
- AI & Systems: apply AI to improve systems, and design systems to improve AI
- Check my publications for details

Assignment 3

- **Intra-domain routing algorithms**

- Design and implement simple versions of link state and distance vector protocols by yourself
- Hands-on experiences on routing protocols: real-world protocols in Cisco routers are just more complicated than the ones you designed and implemented!

- **A preliminary version is online. Not final!**

- Take a look if you want to get a sense of what to do

- **Will add more hints and scaffolding code**

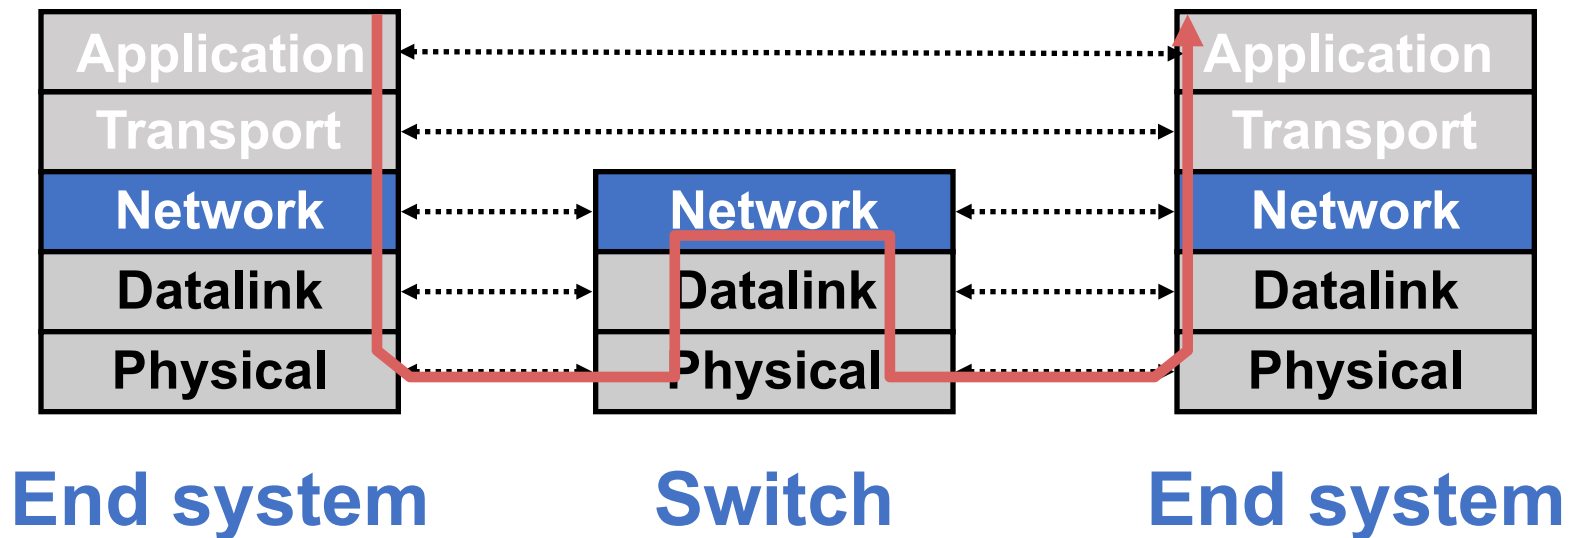
- **Will formally release and discuss it next Monday**

Agenda

- **Network layer recap**
- **Routing fundamentals**

Recap: Network layer

- Present everywhere
- Performs **addressing**, **forwarding**, and **routing**, among other tasks




“Autonomous System (AS)” or “Domain”
Region of a network under a single administrative entity

“End hosts”
“Clients”, “Users”
“End points”


“Route” or “Path”

“Border Routers”

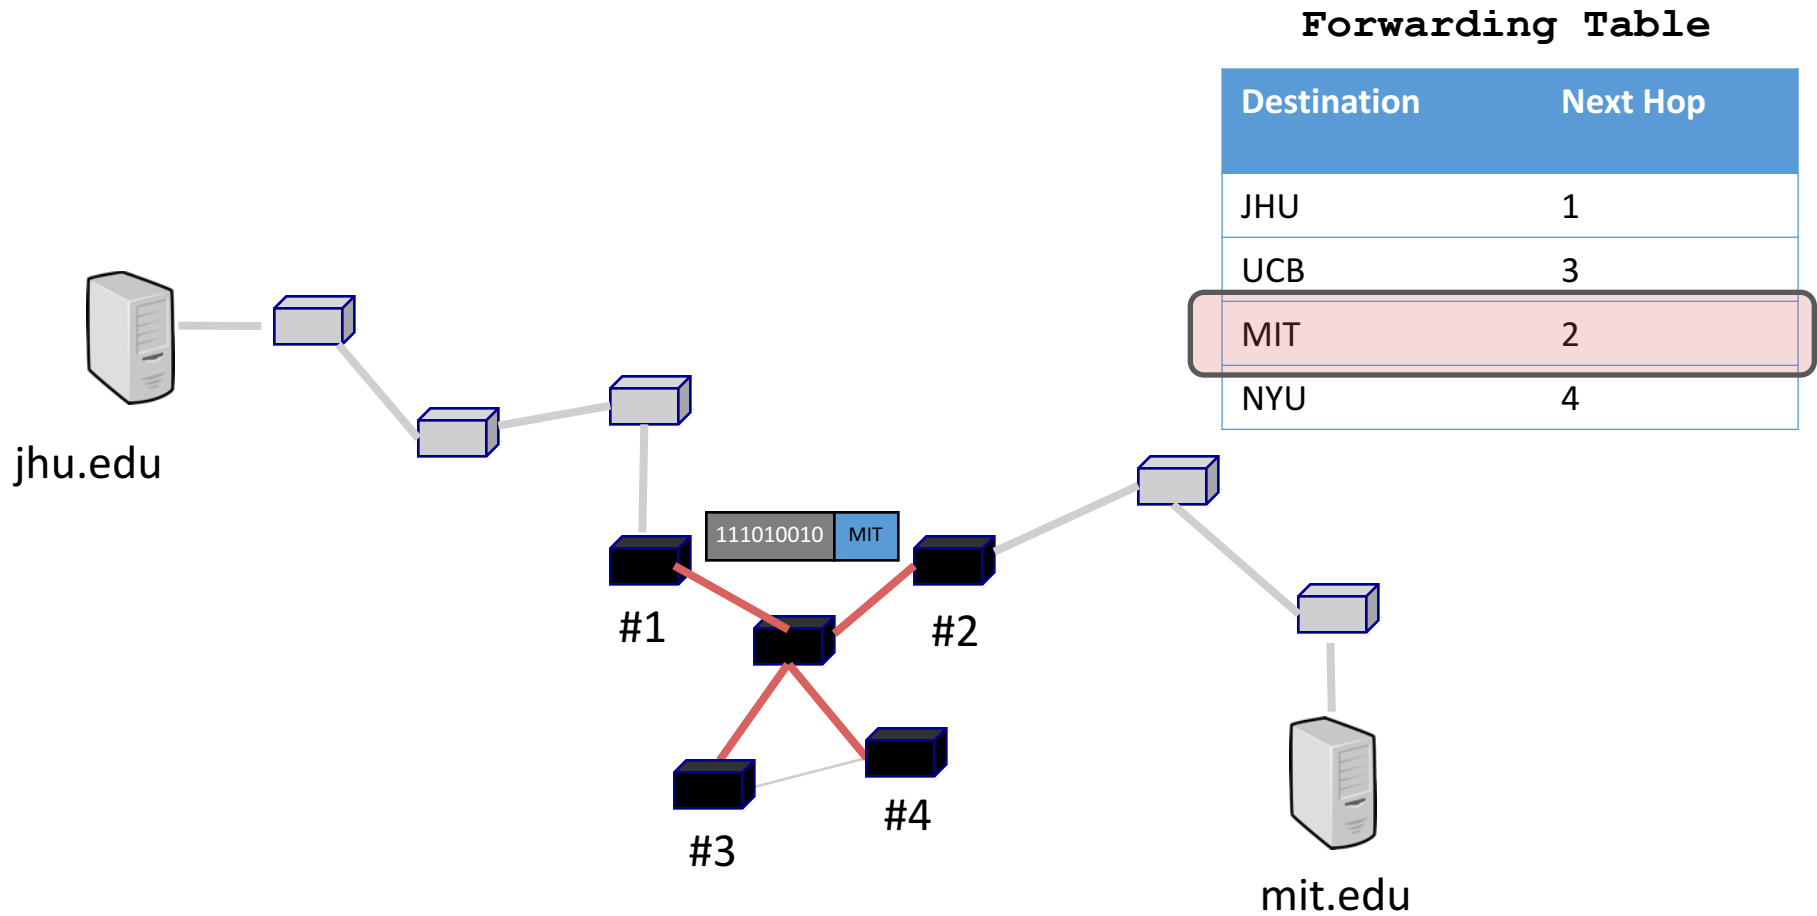
“Interior Routers”



Recap: Forwarding



Recap: Forwarding



Recap: Forwarding

- **Directing a packet to the correct interface so that it progresses to its destination**
 - Local
- **How?**
 - Read address from packet header
 - Search forwarding table

Recap: Routing

- **Setting up network-wide *forwarding tables* to enable end-to-end communication**
 - Global
- **How?**
 - Using different routing protocols

Recap: Forwarding vs. routing

- **Forwarding: “data plane”**
 - Directing one data packet
 - Each router using local routing state
- **Routing: “control plane”**
 - Computing the forwarding tables that guide packets
 - Jointly computed by routers using a distributed algorithm
- **Very different timescales!**

Routing fundamentals

Goal of routing

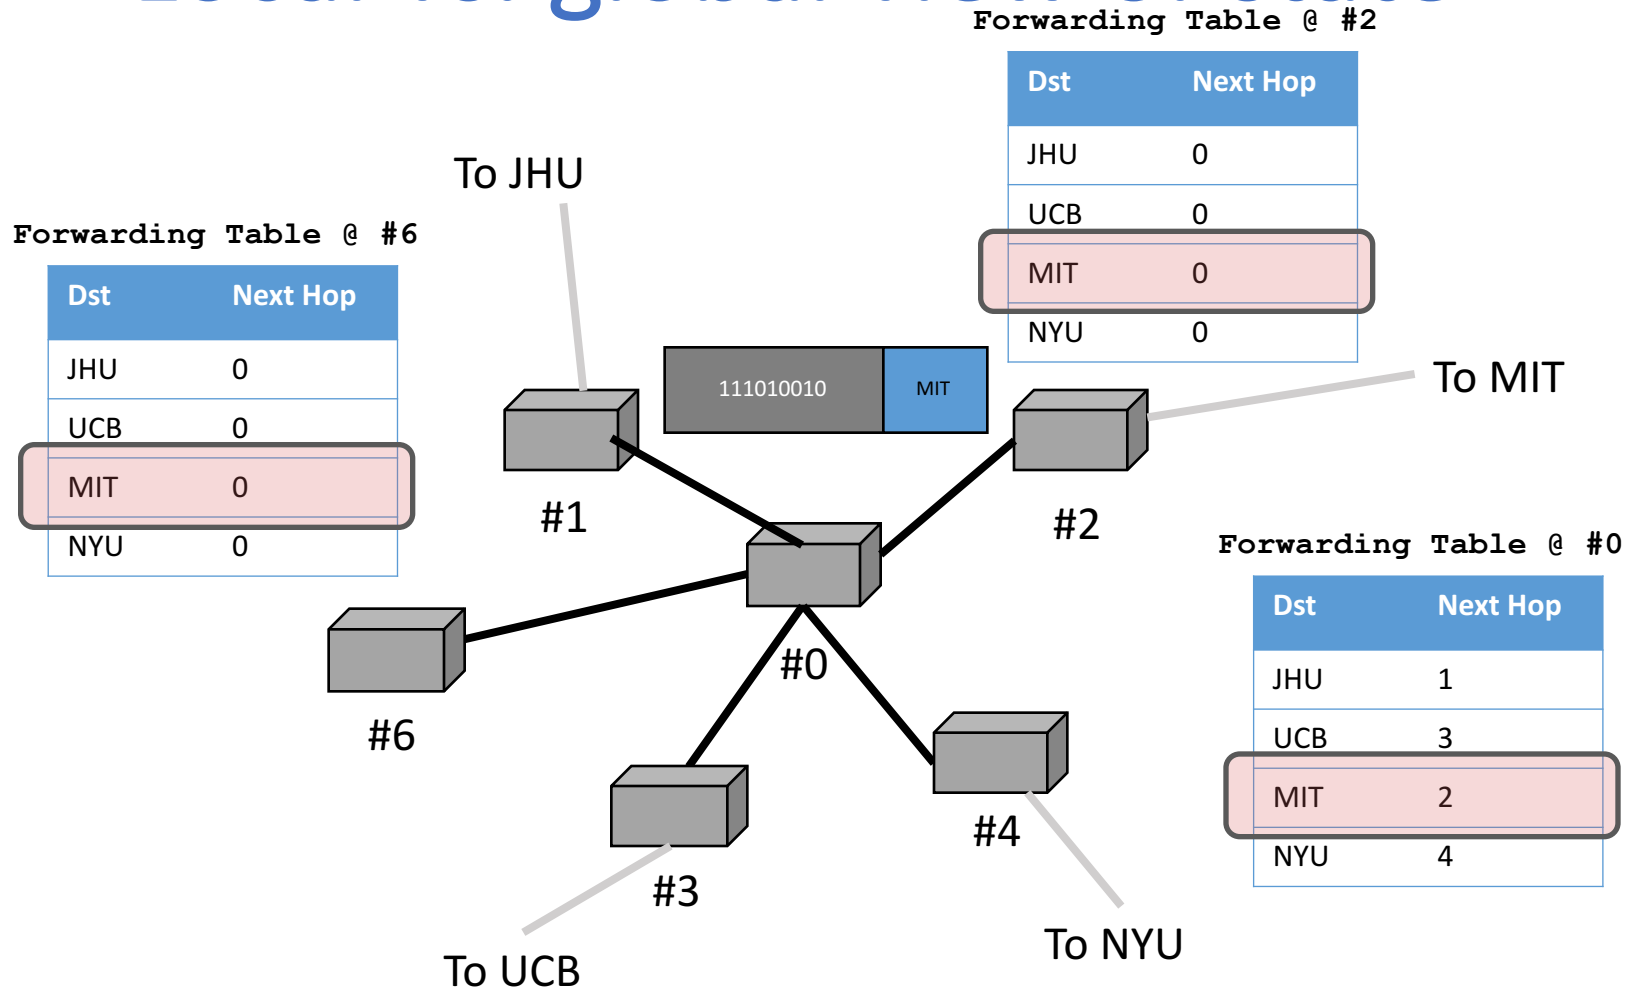
- Find a path to a given destination
- How do we know that the state contained in forwarding tables meets our goal?
 - This is what “**validity**” of routing state tells us
 - [This is non-standard terminology]

Local vs. global view of state

- ***Local* routing state is the forwarding table in a single router**
 - By itself, the state in a single router cannot be evaluated
 - It must be evaluated in terms of the global context

Example:

Local vs. global view of state



Local vs. global view of state

- ***Local* routing state is the forwarding table in a single router**
 - By itself, the state in a single router cannot be evaluated
 - It must be evaluated in terms of the global context
- ***Global* state refers to the collection of forwarding tables in each of the routers**
 - Global state determines which paths packets take
 - (Will discuss later where this routing state comes from)

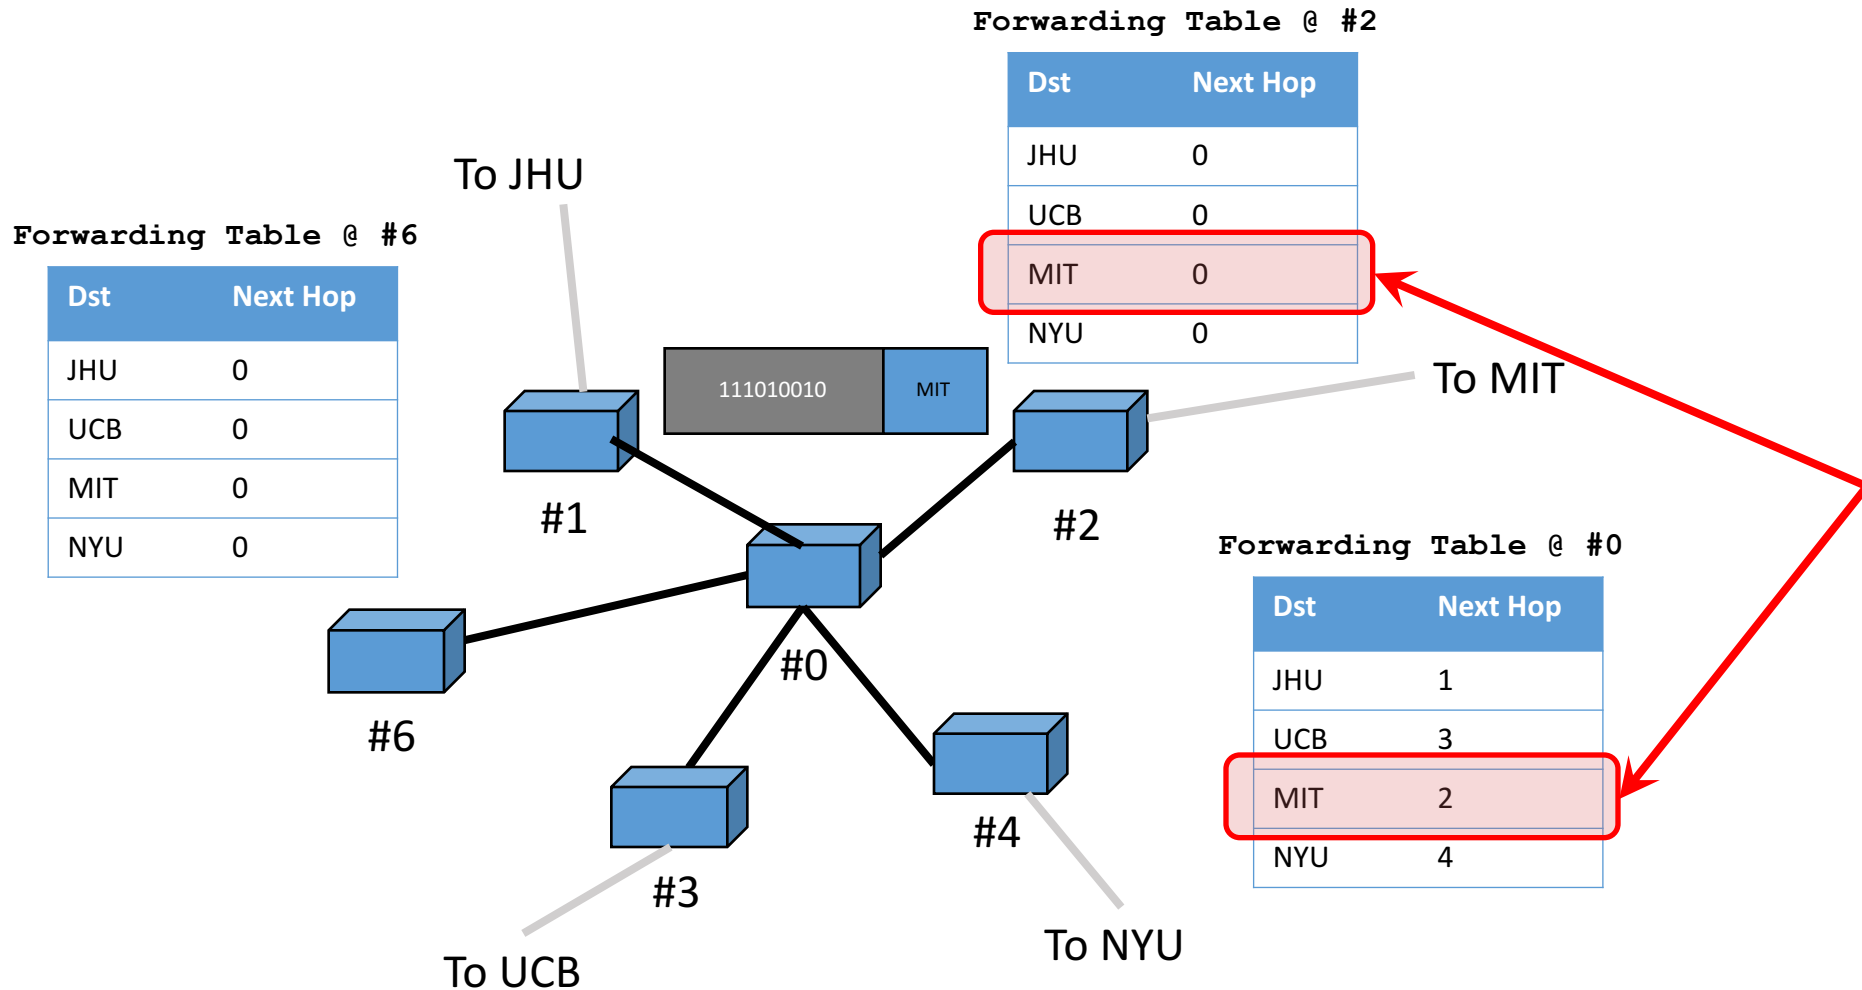
“Valid” routing state

- **Global state is “valid” if it produces forwarding decisions that always deliver packets to their destinations**
- **Goal of routing protocols: compute valid state**
 - How can we tell if routing state is valid?
- **Need a succinct correctness condition for routing**

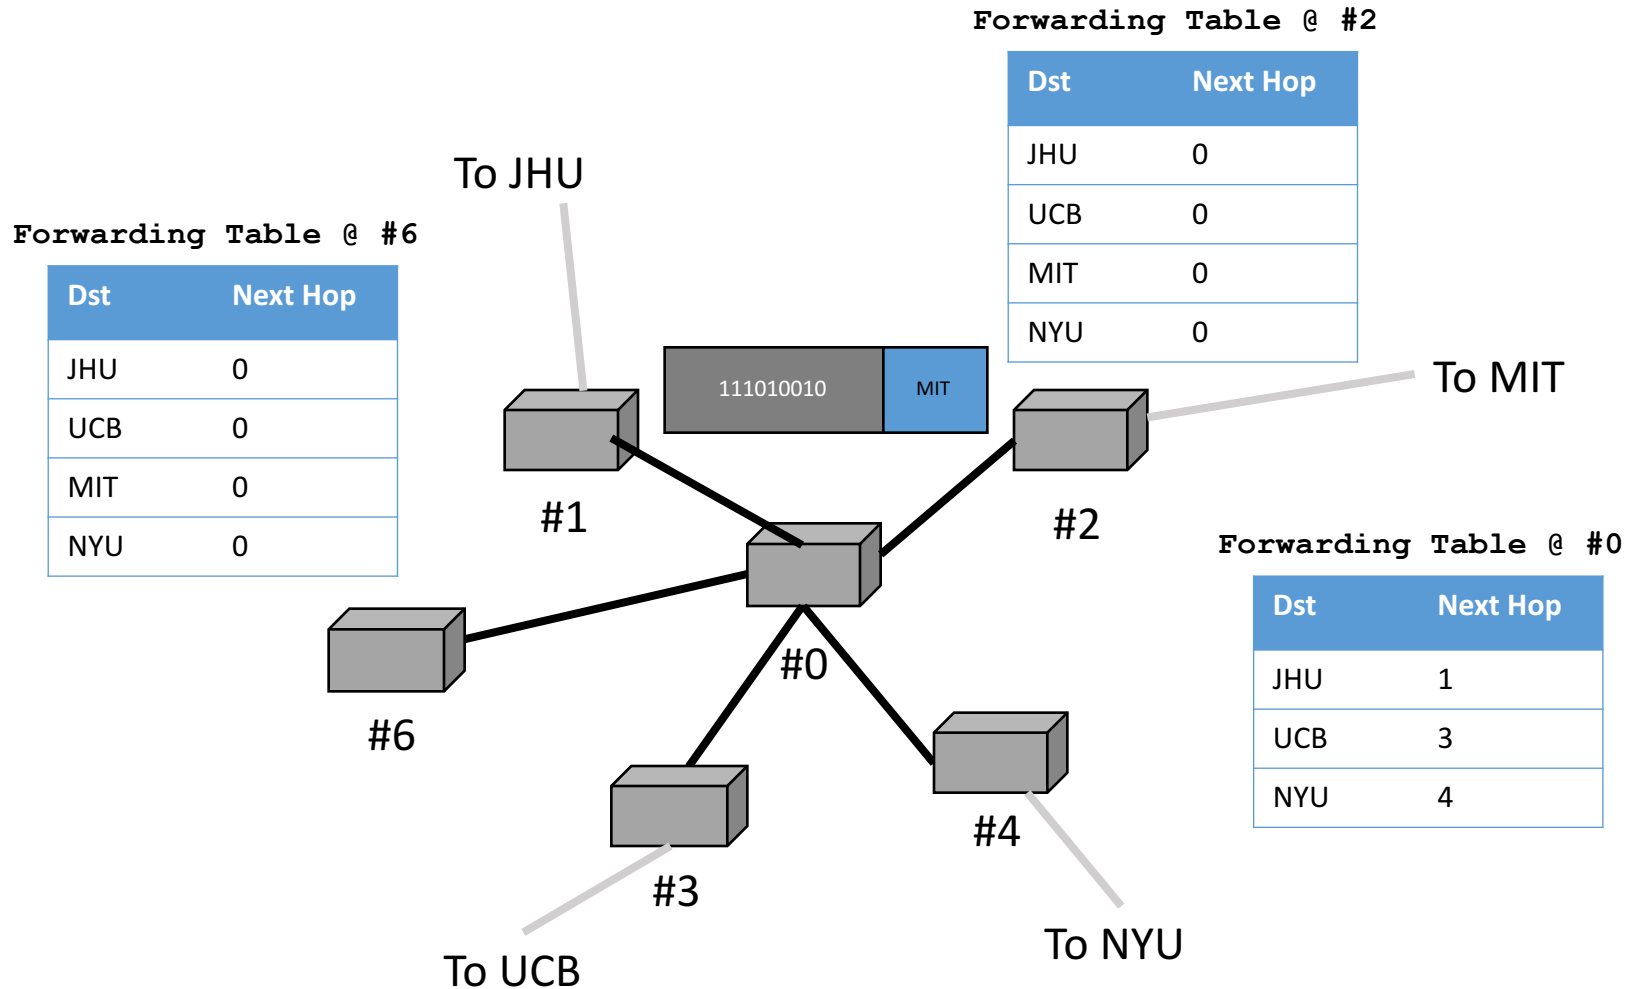
Necessary and sufficient condition

- **Global routing state is valid *if and only if*:**
 - There are no dead ends (other than destination)
 - There are no loops
- **A *dead end* is when there is no outgoing link (next-hop)**
 - A packet arrives, but the forwarding decision does not yield any outgoing link
- **A *loop* is when a packet cycles around the same set of nodes forever**

Loop!



Dead end to MIT @ #0



Necessary and sufficient condition

- **Global routing state is valid *if and only if*:**
 - There are no dead ends (other than destination)
 - There are no loops

Necessary (“only if”)

- **If you run into a dead end before hitting destination,**
 - you'll never reach the destination
- **If you run into a loop,**
 - you'll never reach destination


Sufficient (“if”)

- **Assume there are no dead ends and no loops**
- **Packet must keep wandering, but without repeating**
 - If ever enter same switch from same link, will loop
- **Only a finite number of possible links for it to visit**
 - It cannot keep wandering forever without looping
 - Must eventually hit destination


Checking validity of routing state

- **Focus only on a single destination**
 - Ignore all other routing state
- **Mark outgoing link (“next hop”) with arrow**
 - There is only one at each node
- **Eliminate all links with no arrows**
- **Look at what’s left**


Example 1




Pick destination



Put arrows on outgoing links (to blue dot)




Remove unused links




Leaves spanning tree: Valid

Example 2



Is this valid?

Not valid: Contains loop!



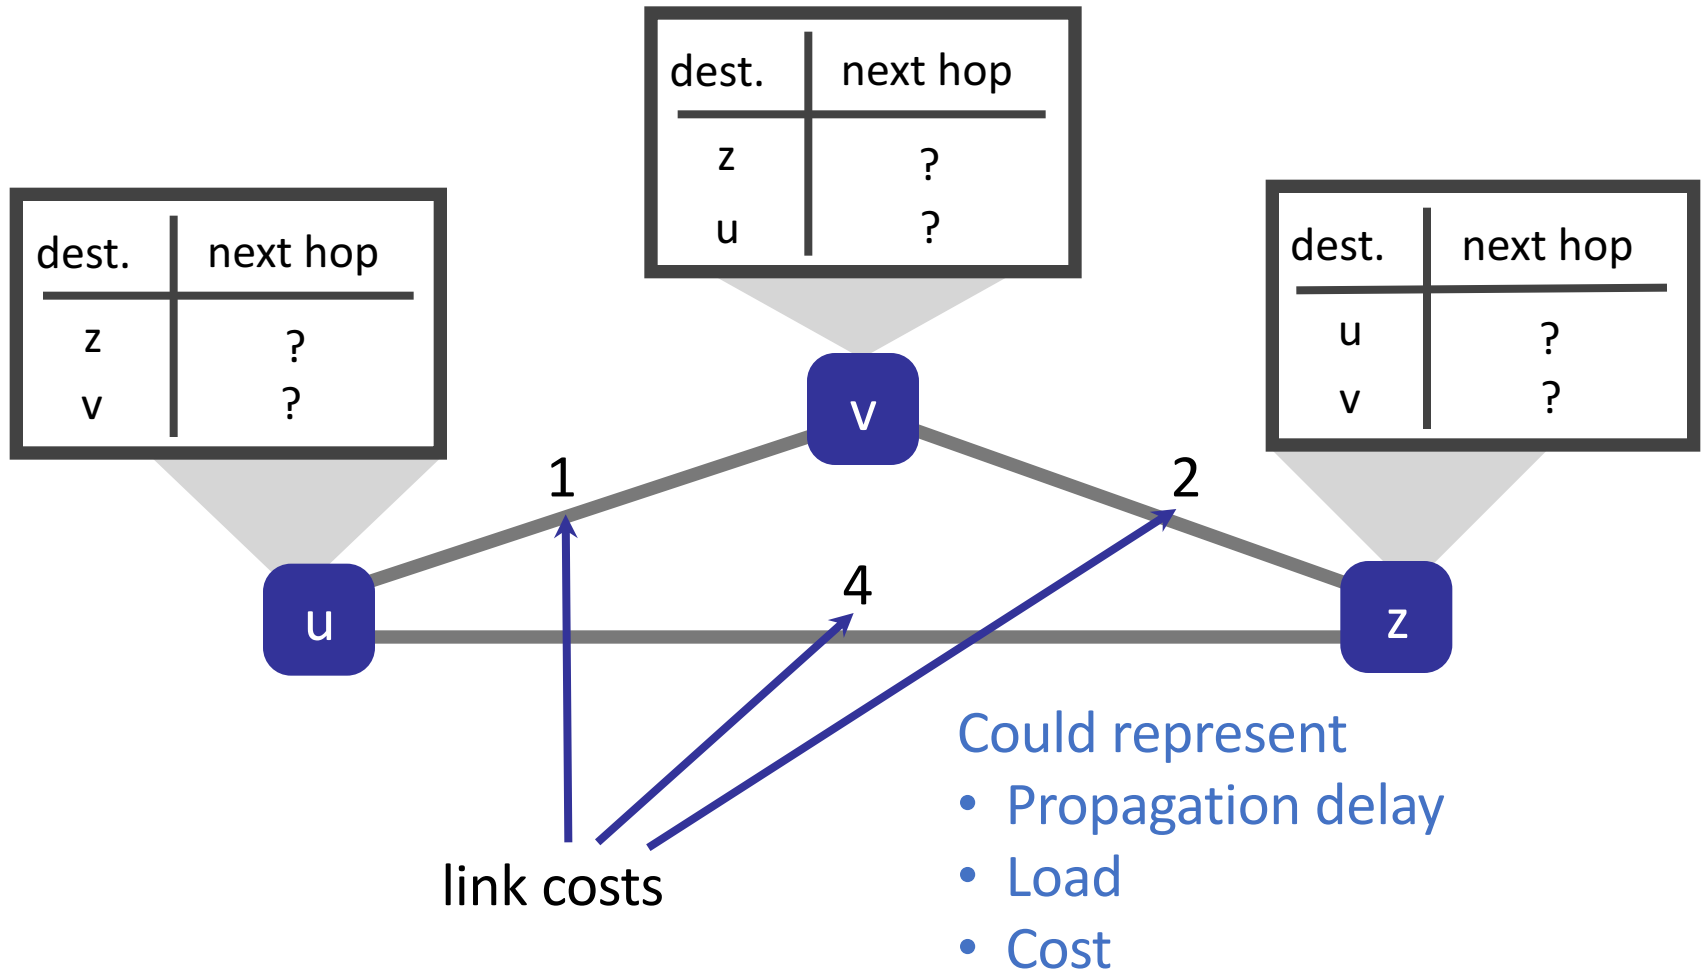
Routing validity

- **Very easy to check validity of routing state for a particular destination**
- **Dead ends are nodes without outgoing arrow**
- **Loops are obvious too**
 - Disconnected from rest of graph


Goal of routing

- **v1:** Find a path to a given destination
- **v2:** Find a *least-cost path* to a given destination

Example



Example



least-cost path from u to z: u v z

least cost path from u to v: u v

Least-cost path routing

- **Given:** router graph & link costs
- **Goal:** find least-cost path
 - From each source router to each destination router

Least-cost routes

- **Least-cost routes provide an easy way to avoid loops**
 - No reasonable cost metric is minimized by traversing a loop
- **Least-cost paths form a spanning tree for each destination rooted at that destination**

Dijkstra's algorithm

- **Network topology, link costs known to all nodes**

- All nodes have same info

- **Computes least-cost paths from one node (“src”) to all other nodes**

- After k iterations, know least-cost path to k destinations

- **Notations**

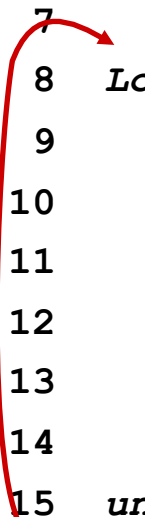
- $c(x,y)$: link cost from x to y ;
 - ∞ if not direct neighbors
- $D(v)$: current value of cost of path from src to dst v
- $p(v)$: predecessor node along path from source to v
- N' : set of nodes whose least-cost path definitively known

Dijkstra's algorithm

```
1  Initialization:  
2     $N' = \{u\}; D(u) = 0$   
3    for all nodes  $v$   
4      if  $v$  adjacent to  $u$   
5        then  $D(v) = c(u, v)$   
6      else  $D(v) = \infty$ 
```

Dijkstra's algorithm

```
1  Initialization:
2    N' = {u}; D(u) = 0
3    for all nodes v
4      if v adjacent to u
5        then D(v) = c(u,v)
6      else D(v) = ∞
7
8  Loop
9    find w not in N' such that D(w) is a minimum
10   add w to N'
11   update D(v) for all v adjacent to w and not in N':
12     D(v) = min( D(v), D(w) + c(w,v) )
13     /* new cost to v is either old cost to v or known
14        least path cost to w plus cost from w to v */
15  until all nodes are in N'
```




Dijkstra's algorithm: Example

Step	N'	D(v) p(v)	D(w) p(w)	D(x) p(x)	D(y) p(y)	D(z) p(z)
0	u	7,u	3,u	5,u	∞	∞
1	uw	6,w		5,u	11,w	∞
2	uw x	6,w			11,w	14,x
3	uw x v				10,v	14,x
4	uw x v y					12,y
5	uw x v y z					

Notes:

- Construct shortest path tree by tracing predecessor nodes
- Ties can exist (can be broken arbitrarily)




Dijkstra's algorithm: Example

*Resulting forwarding table
in u*

destination	link
v	(u, w)
w	(u, w)
x	(u, x)
y	(u, w)
z	(u, w)

*Resulting least-cost tree
from u*



Summary

- **Network layer control plane calculates valid routes and sets up forwarding table**
 - Avoiding loops and dead ends
- **Least-cost routes can be calculated using Dijkstra's algorithm**
- **Next lecture: Routing protocols**

Thanks!
Q&A