

EN.601.414/614

Computer Networks

Routing Algorithms

Xin Jin

Spring 2019 (MW 3:00-4:15pm in Shaffer 301)



Agenda

- **Recap: assignment 2**
- **Recap: least-cost path routing**
- **Link-state routing**
- **Distance-vector routing**
- **Tutorial: assignment 3**

Recap: Assignment 2

- **Primary goal: understand reliable transport**
 - You do not FULLY understand it without implementing it by yourself
- **Secondary goal: it is a lesson for you to learn**
 - Teamwork
 - You and you teammates are on the same boat!
 - Start early, and check progress regularly
 - Software testing
 - You are responsible for the software you write
 - Make sure it passes basic test cases
 - Use the test script of assignment 1 as your start

Recap: Assignment 2

- **Feedback from midterm survey**

- Content: useful and practical

- Reliable transport is used everywhere, and is critical to many applications, e.g., messages, web, remote control, etc.

- Difficulty

- Some students find it too easy, while others find too difficult
 - Tedious things like handling input/output and multi-threading are an integral part of making it work in the real-world
 - Tons of lines of code at Google to make Page-Rank work and build a fast, scalable and reliable search engine
 - We use Python and provide some scaffolding code to abstract away most of the low-level tedious work
 - It is in C, and has no scaffolding code at other schools!

Recap: Assignment 2

- **That being said, changes to assignment 3...**
 - We provide more scaffolding code and the test script (which is the same used for grading)
 - But do try to write the code from scratch, and write your own test script if possible. It is a good learning experience for **your own benefit**.
 - You can also try other languages (C/C++/Java/Go) yourself and build your own test environments
 - Our goal is not to give you a hard time and make you feel bad. We hope you **learn something useful** besides getting A+ (😊) from this course.
 - Finish both distance-vector and link-state protocols to get **bonus points**

Recap: Least-cost path routing

- **Given:** router graph & link costs
- **Goal:** find least-cost path
 - From each source router to each destination router
- **Easy way to avoid loops**
 - No reasonable cost metric is minimized by traversing a loop

Recap:

Dijkstra's algorithm

- **Network topology, link costs known to all nodes**
 - All nodes have same info
- **Each node (“src”) computes least-cost paths to all other nodes**
 - After k iterations, know least-cost path to k destinations


Recap:

Dijkstra's algorithm: Example

Step	N'	D(v) p(v)	D(w) p(w)	D(x) p(x)	D(y) p(y)	D(z) p(z)
0	u	7,u	3,u	5,u	∞	∞
1	uw	6,w		5,u	11,w	∞
2	uw x	6,w			11,w	14,x
3	uw x v			10,v		14,x
4	uw x v y				12,y	
5	uw x v y z					

Notes:

- Construct shortest path tree by tracing predecessor nodes
- Ties can exist (can be broken arbitrarily)



From routing algorithm to protocol

- **Dijkstra's is a local computation!**
 - Computed by a node given complete network graph
- **Possibilities:**
 - **Option#1:** a separate machine runs the algorithm
 - **Option#2:** every router runs the algorithm
- **The Internet currently uses Option#2**

Link-state routing

- **Every router knows its local “link state”**
 - Router u: “(u,v) with cost=2; (u,x) with cost=1”
- **Each router floods its local link state to all other routers in the network**
 - Does so periodically or when its link state changes
- **Every router learns the entire network graph**
 - Each runs Dijkstra’s Shortest-Path First (SPF) algorithm locally to compute forwarding table

Flooding link state

- **Flooding**

- A node sends its link-state info out all of its links
- The next node forwards the info on all of its links except the one the information arrived at

- **When to initiate flooding?**

- Topology change (e.g., link/node failure/recovery)
- Configuration change (e.g., link cost change)
- Periodically
 - To refresh link-state information (soft states)
 - Typically (say) every 30 minutes

Convergence

- **Why flood?**

- To get all the nodes in the network to **converge** to the new topology


- **Upon convergence, all nodes will have **consistent routing information** and can **compute consistent forwarding**:**

- All nodes have the same link-state database
- All nodes forward packets on shortest paths
- The next router on the path forwards to the expected next hop


Convergence delay

- **Time to achieve convergence**
- **Sources of convergence delay**
 - Time to detect failure
 - Time to flood link-state information
 - Time to re-compute forwarding tables
- **What happens if it takes too long to converge?**

Loop from convergence delay



u and w think that the path to y goes through v



v thinks that the path to y goes through w

Performance during convergence period

- Looping packets
- Lost packets **due to black holes**
- Out-of-order packets **reaching the destination**

Link-state routing

- **Scalability?**

- $O(NE)$ messages
- $O(N^2)$ computation time
- $O(\text{Network diameter})$ convergence delay
- $O(N)$ entries in forwarding table

Link-state routing protocols

- **OSPF**: Open Shortest Path First
- **IS-IS**: Intermediate System to Intermediate System
 - Similar to OSPF

OSPF:

Open Shortest-Path First

- **Open:** publicly available
- **Uses link-state algorithm**
 - Link-state packet dissemination
 - Topology map at each node
 - Route computation using Dijkstra's algorithm
- **Router floods OSPF link-state advertisements to all other routers in entire AS**
 - Carried in OSPF messages directly over IP (rather than TCP or UDP)
 - Requires reliable transmission

Distance-vector protocol

- **Link-state routing protocol**

- Each node **broadcasts** its **local** information

- **Distance-vector routing protocol**

- The opposite (sort of)

- Each node **tells its neighbors** about its **global** view

Bellman-Ford equation

- **Let**

➤ $d_x(y) :=$ cost of least-cost path from x to y

- **Then**


➤ $d_x(y) = \min_v \{ c(x, v) + d_v(y) \}$

\min taken over all neighbors v of x

cost to neighbor v

cost from neighbor v to destination y

Bellman-Ford example



$$d_x(z) = 9, d_w(z) = 9, d_v(z) = 6$$

B-F equation says:

$$\begin{aligned} d_u(z) &= \min \{c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z), \\ &\quad c(u,v) + d_v(z) \} \\ &= \min \{5 + 9, \\ &\quad 3 + 9, \\ &\quad 7 + 6\} = 12 \end{aligned}$$

Neighbor achieving the minimum (w) is next hop in shortest path, used in forwarding table

Distance vector algorithm

- $D_x(y)$ is the estimate of least cost from x to y
 - x maintains its own distance vector $D_x = [D_x(y): y \in N]$
- **Node x :**
 - Knows cost to each neighbor v : $c(x,v)$
 - Maintains its neighbors' distance vectors
 - For each neighbor v , x has $D_v = [D_v(y): y \in N]$

Distance vector algorithm

- From time-to-time, each node sends its own distance vector estimate to neighbors
- When x receives new DV estimate from neighbor, it updates its own DV using B-F equation
 - $D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\}$ for each node $y \in N$
- Eventually, the estimate $D_x(y)$ may converge to the actual least cost $d_x(y)$

Distance vector algorithm

- **Iterative, asynchronous**

- Local iterations caused by

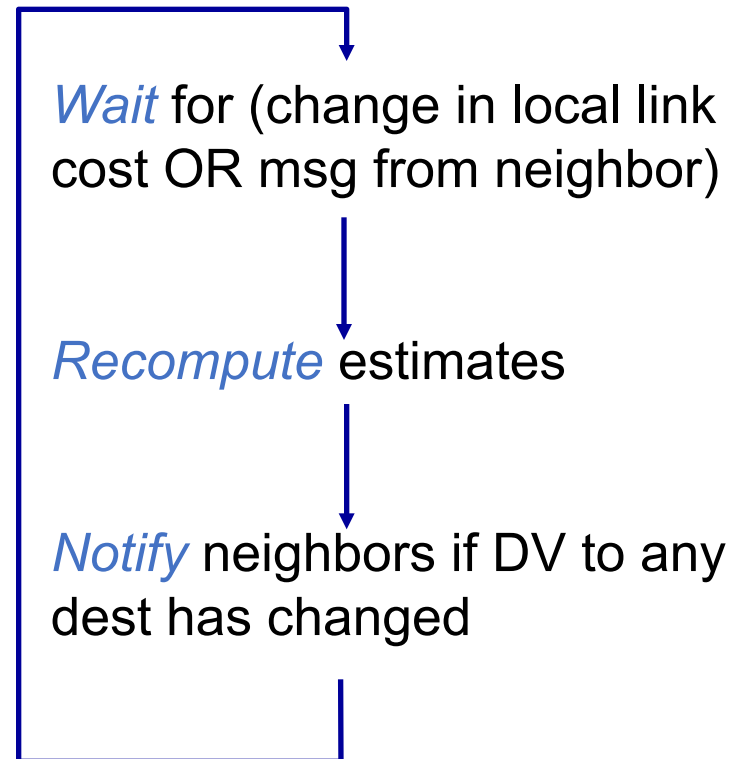
- Local link cost change
- DV update message from neighbor

- **Distributed**

- Each node notifies neighbors only when its DV changes


- Neighbors then notify their neighbors if necessary

@each node:



Example

	x	y	z
y	2	0	1
z	3	1	0




	x	y	z
y	2	0	1
z	3	1	0

Example

	x	y	z
y	4	0	1
z	3	1	0

routing loop!




	x	y	z
y	2	0	1
z	3	1	0

Example

	x	y	z
y	4	0	1
z	3	1	0

routing loop!




	x	y	z
y	4	0	1
z	3	1	0

Example

	x	y	z
y	4	0	1
z	3	1	0

routing loop!




	x	y	z
y	4	0	1
z	5	1	0

Example

	x	y	z
y	4	0	1
z	5	1	0

routing loop!




	x	y	z
y	4	0	1
z	5	1	0

Example

	x	y	z
y	6	0	1
z	5	1	0

routing loop!




	x	y	z
y	4	0	1
z	5	1	0

Example

	x	y	z
y	6	0	1
z	5	1	0

routing loop!




	x	y	z
y	6	0	1
z	5	1	0

Example

	x	y	z
y	6	0	1
z	5	1	0

routing loop!




	x	y	z
y	6	0	1
z	7	1	0

Example

	x	y	z
y	6	0	1
z	7	1	0

routing loop!




	x	y	z
y	6	0	1
z	7	1	0

Example

	x	y	z
y	8	0	1
z	7	1	0

routing loop!




	x	y	z
y	6	0	1
z	7	1	0

Example

	x	y	z
y	8	0	1
z	7	1	0

routing loop!




	x	y	z
y	8	0	1
z	7	1	0

Example

	x	y	z
y	8	0	1
z	7	1	0

routing loop!



Count-to-infinity scenario

	x	y	z
y	8	0	1
z	7	1	0


Problems with Bellman-Ford

- **Routing loops**
 - z routes through y, y routes through x
 - y loses connectivity to x
 - y decides to route through z
- **Can take a very long time to resolve**
 - Count-to-infinity scenario

Poisoned reverse


- One **heuristic** to avoid count-to-infinity
 - If z routes to x through y,
 - z advertises to y that its cost to x is infinite
 - y never decides to route to x through z

	x	y	z
y	2	0	1
z	∞	1	0




	x	y	z
y	2	0	1
z	3	1	0

	x	y	z
y	∞	0	1
z	∞	1	0




	x	y	z
y	2	0	1
z	3	1	0

	x	y	z
y	∞	0	1
z	∞	1	0




	x	y	z
y	∞	0	1
z	3	1	0

	x	y	z
y	∞	0	1
z	∞	1	0




	x	y	z
y	∞	0	1
z	7	1	0

	x	y	z
y	∞	0	1
z	7	1	0



	x	y	z
y	∞	0	1
z	7	1	0

	x	y	z
y	8	0	1
z	7	1	0



	x	y	z
y	∞	0	1
z	7	1	0

Poisoned reverse

- **One heuristic to avoid count-to-infinity**
 - If z routes to x through y,
 - z advertises to y that its cost to x is infinite
 - y never decides to route to x through z
- **Not guaranteed (why?)**
- **Loop-free routing examples include**
 - Path vector
 - Source tracing

Distance-vector routing

- **Scalability?**

- Requires fewer messages than Link-State
- $O(N)$ update time on arrival of a new DV from neighbor
- $O(\text{network diameter})$ convergence time
- $O(N)$ entries in forwarding table

- **RIP is a protocol that implements DV (IETF RFC 2080)**

Comparison of LS and DV routing

Messaging complexity

- **LS:** with N nodes, E links, $O(NE)$ messages sent
- **DV:** exchange between neighbors only

Speed of convergence

- **LS:** relatively fast
- **DV:** convergence time varies
 - Count-to-infinity problem

Robustness: what happens if router malfunctions?

- **LS:**
 - Node can advertise incorrect **link** cost
 - Each node computes its *own* table
- **DV:**
 - Node can advertise incorrect **path** cost
 - Each node's table used by others (error propagates)

Similarities between LS and DV routing

- **Both are shortest-path based routing**
 - Minimizing cost metric (link weights) a common optimization goal
 - Routers share a common view as to what makes a path “good” and how to measure the “goodness” of a path
- **Due to shared goal, commonly used inside an organization**
 - RIP and OSPF are mostly used for **intra**-domain routing

Tutorial: Assignment 3

- **Intra-domain routing algorithms**

- Design and implement simple versions of link state and distance vector protocols by yourself
- Hands-on experiences on routing protocols: real-world protocols in Cisco routers are just more complicated than the ones you designed and implemented!
- Use NetworkX to compute shortest path if you do not want to implement it yourself

- **Implement either link-state or distance-vector**

- Implementing both gives you a bonus of 2 points

Summary

- **Intra-AS routing**
 - Link-state routing
 - Distance-vector routing
- **Next class: Midterm re-cap**
- **Next Monday: Inter-AS routing**

Thanks!
Q&A