

EN.601.414/614

Computer Networks

Midterm Review

Xin Jin

Fall 2020 (TuTh 1:30-2:45pm on Zoom)



<https://github.com/xinjin/course-net>

Midterm Exam

- **Time: 75 minutes on October 15 (this Thursday)**
- **Location: Take-home**
- **Form: Open-book**
 - Can use slides for reference
 - Can use a calculator
 - Anything else is prohibited
- **Senior Option**
 - Your final exam score will be the same as your midterm exam score

This review

- **Walk through what you're expected to know at this point: key topics, important aspects of each**
- **Not covered in review does NOT imply you don't need to know it**
 - But if it's covered today, you should know it
- **Summarize, not explain**
 - Stop me when you want to discuss something further!

Topics

- **Basics (lectures 1–3)**
- **Application layer (lectures 4, 5)**
 - HTTP, DNS, and CDN
- **Transport layer (lectures 6–9)**
 - UDP vs. TCP
 - TCP details: reliability and flow control
 - TCP congestion control: general concepts only
- **Network layer (lectures 10, 11)**
 - Data plane

Basic concepts

- **You should know:**

- Packet vs. circuit switching
- Statistical multiplexing
- Link characteristics
- Packet delays

How are network resources shared?

- **Two approaches**

- Reservations → circuit switching
- On-demand → packet switching

Two approaches to sharing

- **Packet switching**


- Network resources consumed on demand per-packet
- Admission control: per packet

- **Circuit switching**

- Network resources reserved a priori at “connection” initiation
- Admission control: per connection

Circuit switching

1. **src** sends reservation request to **dst**
2. Switches create circuit *after* admission control
3. **src** sends data
4. **src** sends teardown request



Packet switching

- Data is sent as chunks of formatted bits (Packets)
- Packets consist of a “header” and “payload”
- Switches “forward” packets based on their headers
- Each packet travels **independently**
- No link resources are reserved in advance

Statistical multiplexing

- **Allowing more demands than the network can handle**
 - Hoping that not all demands are required at the same time
 - Good for bursty traffic (average \ll peak demand)
 - Packet switching exploits statistical multiplexing better than circuit switching

Performance metrics

- **Delay**
- **Loss**
- **Throughput**

A network link



- **Link bandwidth**
 - Number of bits sent/received per unit time (bits/sec or bps)
- **Propagation delay**
 - Time for one bit to move through the link (seconds)
- **Bandwidth-Delay Product (BDP)**
 - Number of bits “in flight” at any time
- **$\text{BDP} = \text{bandwidth} \times \text{propagation delay}$**

Delay

- **Consists of four components**

➤ Transmission delay

➤ Propagation delay



due to link properties


➤ Queuing delay

➤ Processing delay




due to traffic mix and
switch internals

End-to-end delay



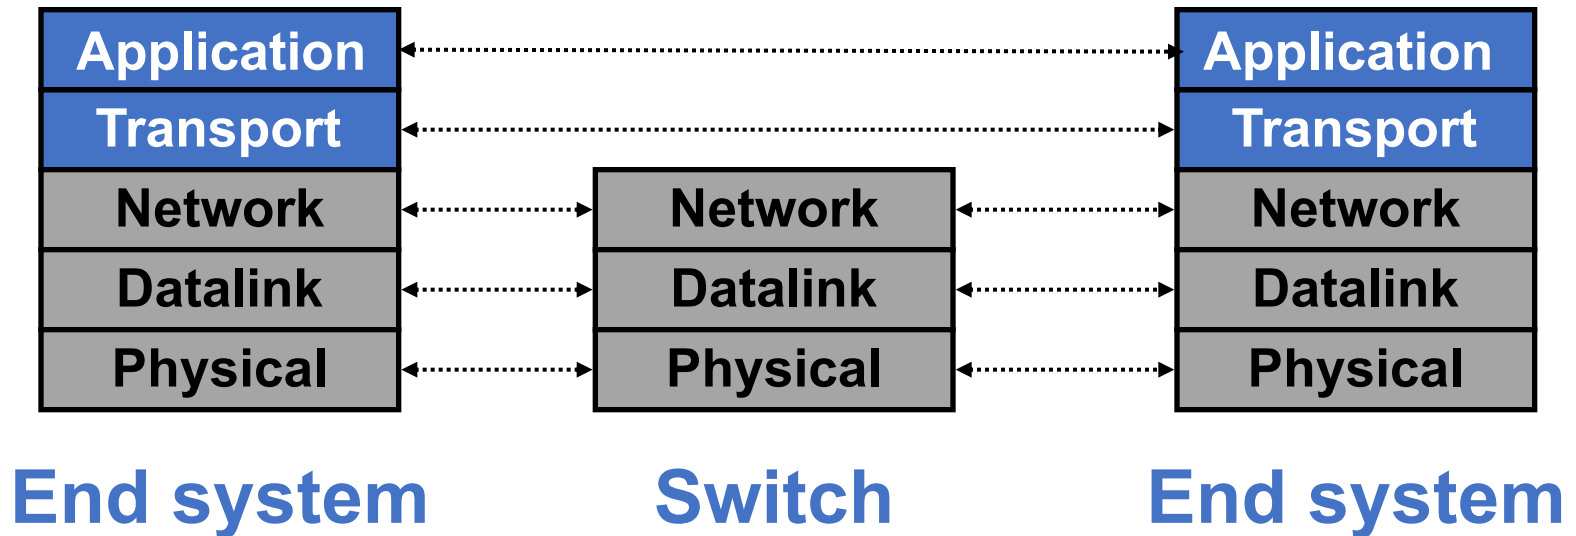
OSI layers

- **OSI stands for Open Systems Interconnection model**
 - Developed by the ISO
- **Session and presentation layers are often implemented as part of the application layer**

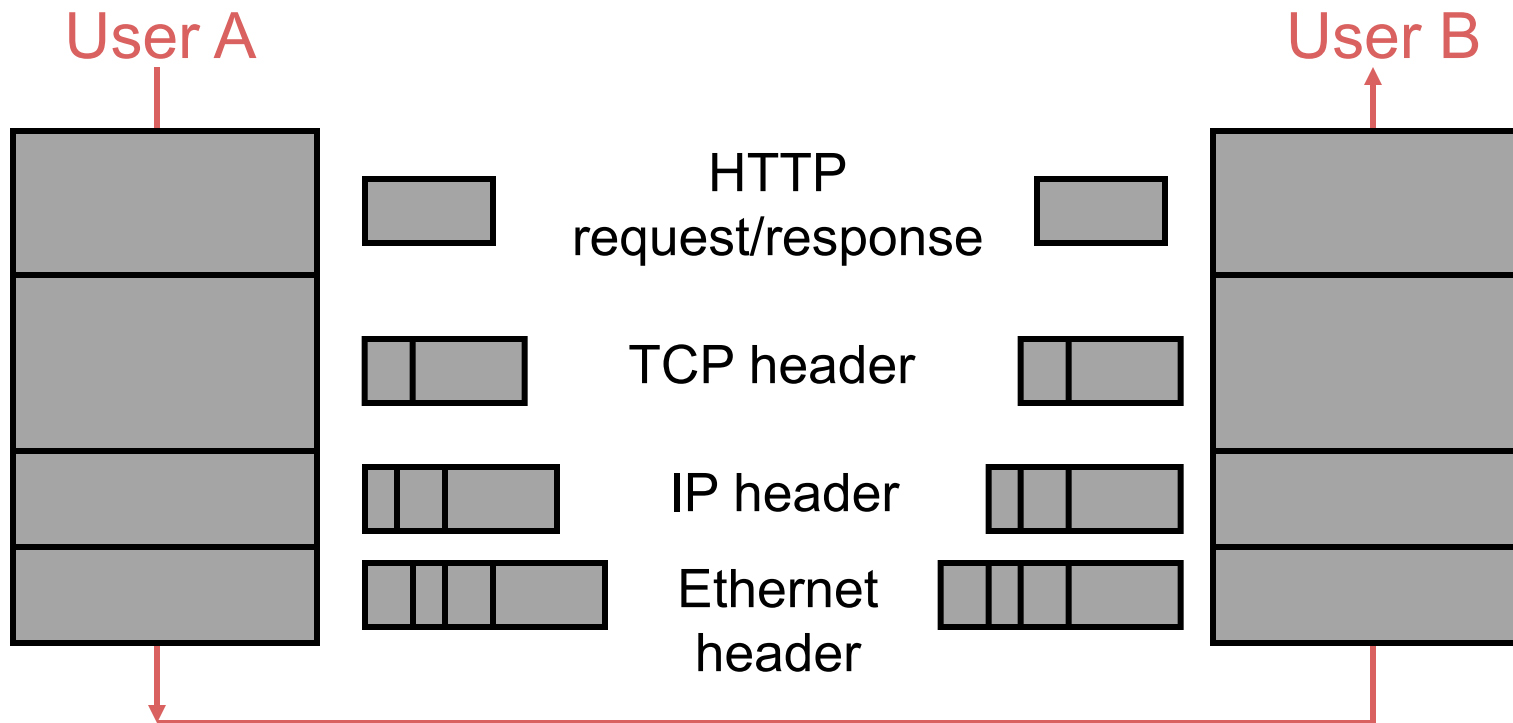


Layers in practice

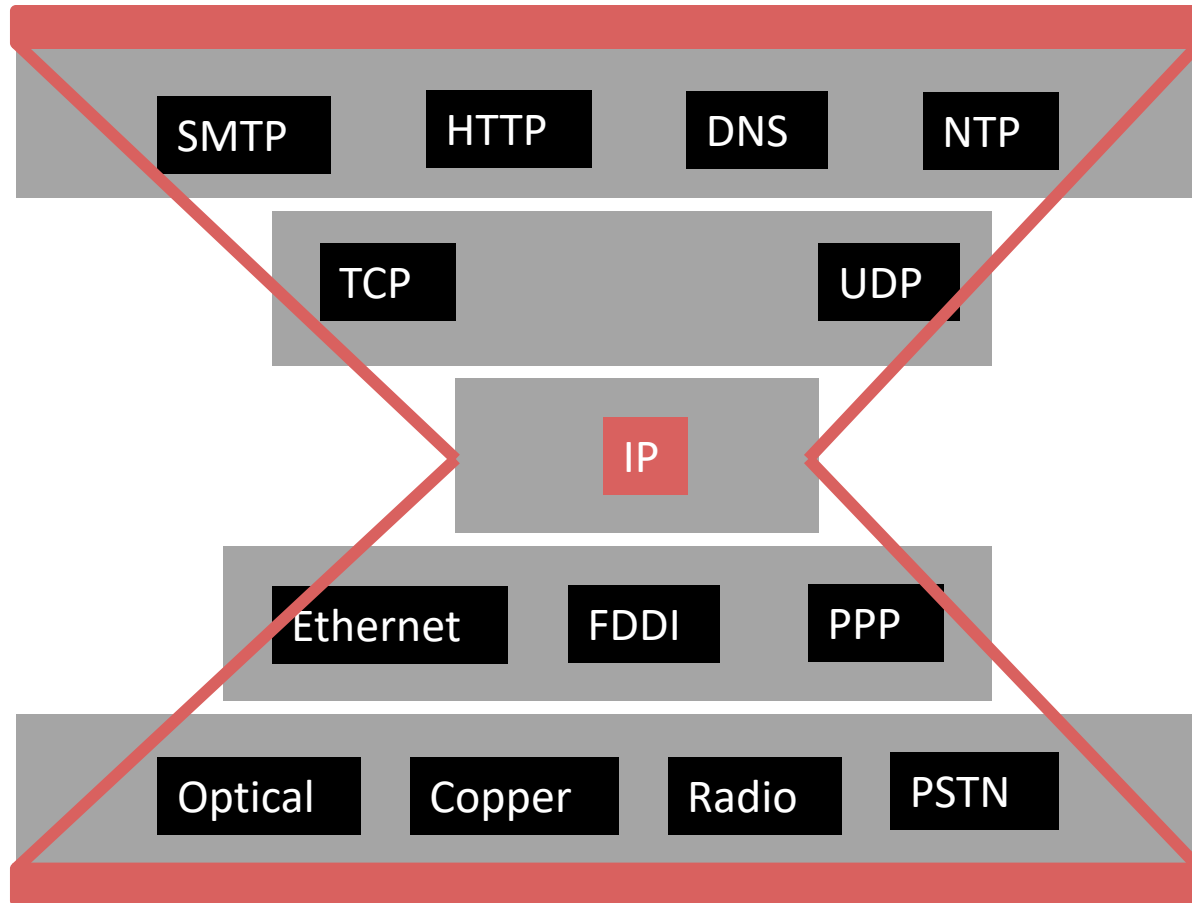
- Lower three layers implemented everywhere
- Top two layers implemented only at hosts



Layer encapsulation: Protocol headers



IP is the narrow waist of the layering hourglass



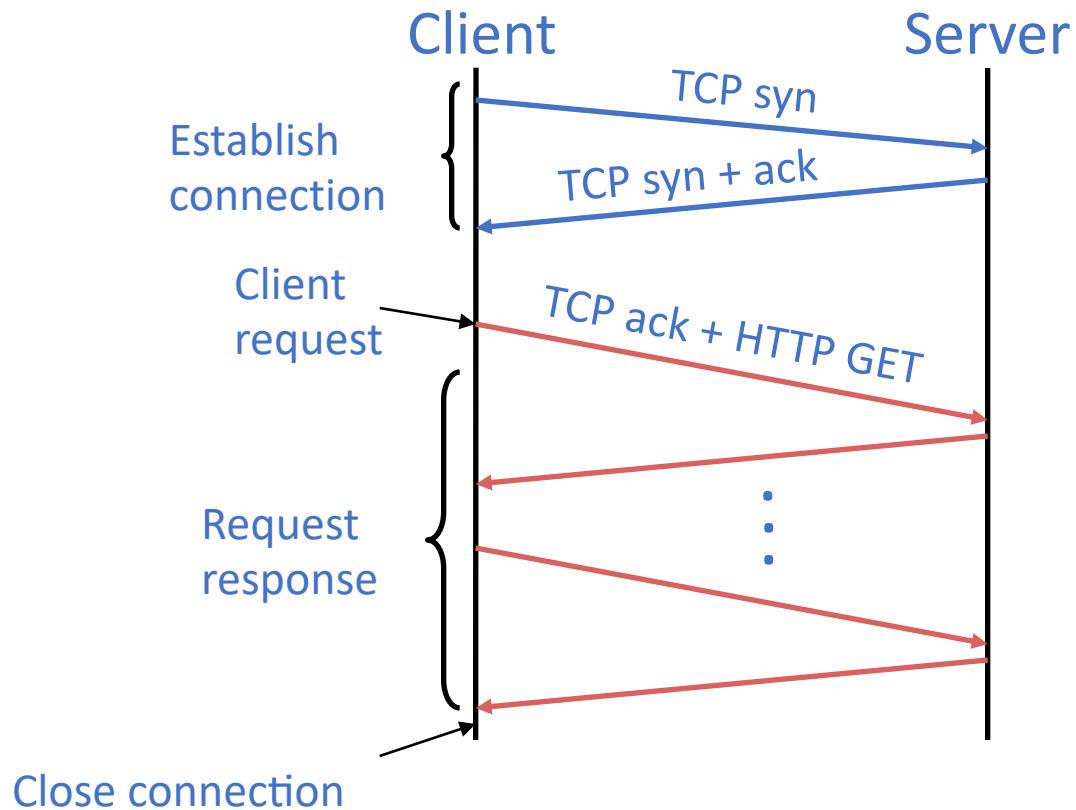
Topics

- Basics (lectures 1–3)
- **Application layer (lectures 4, 5)**
 - HTTP, DNS, and CDN
- **Transport layer (lectures 6–9)**
 - UDP vs. TCP
 - TCP details: reliability and flow control
 - TCP congestion control: general concepts only
- **Network layer (lectures 10, 11)**
 - Data plane

Hyper Text Transfer Protocol (HTTP)

- **Client-server architecture**
 - Server is “always on” and “well known”
 - Clients initiate contact to server
- **Synchronous request/reply protocol**
 - Runs over TCP, Port 80
- **Stateless**
- **ASCII format**
 - Before HTTP/2

Steps in HTTP request/response




Object request response time

- **RTT (round-trip time)**

- Time for a small packet to travel from client to server and back

- **Response time**

- 1 RTT for TCP setup
- 1 RTT for HTTP request and first few bytes
- Transmission time
- **Total** = 2RTT + Transmission Time



Improving HTTP performance

- **Optimizing connections using three “P”s**
 - Persistent connections
 - Parallel/concurrent connections
 - Pipelined transfers over the same connection
- **Caching**
 - Forward proxy: close to clients
 - Reverse proxy: close to servers
- **Replication**


Content Distribution Networks (CDN)

- **Caching and replication as a service**
- **Combination of caching and replication**
 - **Pull**: Direct result of clients' requests (caching)
 - **Push**: Expectation of high access rate (replication)


Hierarchies in the DNS

- **Three intertwined hierarchies**
 - Hierarchical namespace
 - As opposed to original flat namespace
 - Hierarchically administered
 - As opposed to centralized
 - (Distributed) hierarchy of servers
 - As opposed to centralized storage


Name resolution




Name resolution




Name resolution




Name resolution



Name resolution: Recursive



Name resolution: Iterative



DNS caching

- **Performing all these queries takes time**
 - Up to 1-second latency before starting download
- **Caching can greatly reduce overhead**
 - The top-level servers very rarely change
 - Popular sites (e.g., www.cnn.com) visited often
 - Local DNS server often has the information cached
- **How DNS caching works**
 - DNS servers cache responses to queries
 - Responses include a “time to live” (TTL) field
 - Server deletes cached entry after TTL expires

Topics

- Basics (lectures 1–3)
- Application layer (lectures 4, 5)
 - HTTP, DNS, and CDN
- **Transport layer (lectures 6–9)**
 - UDP vs. TCP
 - TCP details: reliability and flow control
 - TCP congestion control: general concepts only
- **Network layer (lectures 10, 11)**
 - Data plane

Role of the transport layer

- **(1) Communication between application processes**
 - Mux and demux from/to application processes
 - Implemented using ports
- **(2) Provide common end-to-end services for app layer**
 - Reliable, in-order data delivery
 - Well-paced data delivery

UDP vs. TCP

➤ Both UDP and TCP perform mux/demux via ports

	UDP	TCP
Data abstraction	Packets (datagrams)	Stream of bytes of arbitrary length
Service	Best-effort (same as IP)	<ul style="list-style-type: none">•Reliability•In-order delivery•Congestion control•Flow control

Reliable transport: General concepts

- **Checksums (for error detection)**
- **Timers (for loss detection)**
- **Acknowledgments (feedback from receiver)**
 - Cumulative: “received everything up to X”
 - Selective: “received X”
- **Sequence no (detect duplicates, accounting)**
- **Sliding windows (for efficiency)**

You should know:

- what these concepts are
- why they exist
- how TCP uses them

Designing a reliable transport protocol

- **Stop and wait is correct but inefficient**
 - Works packet by packet (of size DATA)
 - Throughput is $(\text{DATA} / \text{RTT})$
- **Sliding window: use pipelining to increase throughput**
 - n packets at a time results in higher throughput
 - $\text{MIN}(n * \text{DATA} / \text{RTT}, \text{Link Bandwidth})$

The TCP abstraction

- **TCP delivers a reliable, in-order, byte stream**
- **Reliable:** TCP resends lost packets (recursively)
 - Until it gives up and shuts down connection
- **In-order:** TCP only hands consecutive chunks of data to application
- **Byte stream:** TCP assumes there is an incoming stream of data, and attempts to deliver it to app

Things to know about TCP

- **How TCP achieves reliability**
 - **RTT estimation**
 - **Connection establishment/teardown**
 - **Flow Control**
 - **Congestion Control**
-
- **For each, know how the functionality is implemented and why it is needed**

Reliability

- **Having TCP take care of it simplifies application development**
- **How**
 - Checksums and timers (for error and loss detection)
 - Fast retransmit (to detect faster-than-timeout loss)
 - Cumulative ACKs (receiver feedback: what's lost?)
 - Sliding windows (for efficiency)
 - Buffers at sender (hold packets until ACKs arrive)
 - Buffers at receiver (to reorder packets before delivery to application)

RTT estimation


- **TCP uses timeouts to retransmit packets**
 - But RTT may vary (significantly!) for different reasons and on different timescales
 - due to temporary congestion
 - due to long-lived congestion
 - due to a change in routing paths
- **An incorrect RTT estimate might introduce spurious retransmissions or overly long delays**
- **Proposed solutions use EWMA, incorporate deviations**

Establishing a TCP connection

- **Three-way handshake to establish connection**

- Host A sends a SYN (open; “synchronize sequence numbers”) to host B
- Host B returns a SYN acknowledgment (SYN ACK)
- Host A sends an ACK to acknowledge the SYN ACK

- **Three-way handshake to terminate (normal operation)**



Flow control

- **Why?**

- TCP at the receiver must buffer a packet until all packets before it (in byte-order) have arrived and the receiving application has consumed available bytes
- Hence, receiver advances its window when the receiving application consumes data
- Sender advances its window when new data ACK'd
- Risk of sender over-running the receiver's buffers

- **How?**

- “Advertised Window” field in TCP header

Congestion control

- **Why?**

- Because the network itself can be the bottleneck
- Should make efficient use of available network capacity
 - While sharing available capacity fairly with other flows
 - And adapting to changes in available capacity

- **How?**

- Dynamically adapts the size of the sending window

Put together

- **Flow Control**

- Restrict window to RWND to make sure that the receiver isn't overwhelmed

- **Congestion Control**

- Restrict window to CWND to make sure that the network isn't overwhelmed

- **Together**

- Restrict window to $\min\{\text{RWND}, \text{CWND}\}$ to make sure that neither the receiver nor the network are overwhelmed

CC implementation

- **States at sender**

- **CWND** (initialized to a small constant)
- **ssthresh** (initialized to a large constant)
- **dupACKcount** and **timer**

- **Events**

- **ACK** (new data)
- **dupACK** (duplicate ACK for old data)
- **Timeout**

Event: ACK (new data)

- If **CWND** < **sssthresh**

➤ **CWND** += 1

- *CWND* packets per RTT
- Hence, after one RTT with no drops:
 $CWND = 2 \times CWND$

Event: ACK (new data)

- If **CWND** < **ssthresh**

- **CWND** += 1

Slow start phase

- **Else**

- **CWND** = **CWND** + 1/**CWND**

Congestion avoidance phase

- *CWND* packets per RTT
- Hence, after one RTT with no drops:
 $CWND = CWND + 1$

Event: Timeout

- **On Timeout**


- $ssthresh \leftarrow CWND/2$

- $CWND \leftarrow 1$

Event: dupACK

- **dupACKcount ++**
- **If dupACKcount = 3 /* fast retransmit */**
 - **ssthresh = CWND/2**
 - **CWND = CWND/2**

Example



Slow-start restart: Go back to $\text{CWND} = 1 \text{ MSS}$, but take advantage of knowing the previous value of CWND

TCP flavors

- **TCP-Tahoe**

- $CWND = 1$ on 3 dupACKs

- **TCP-Reno**

- $CWND = 1$ on timeout


- $CWND = CWND/2$ on 3 dupACKs

- **TCP-newReno**

- TCP-Reno + improved fast recovery


- **TCP-SACK**

- Incorporates selective acknowledgements




Our default assumption

A simple model for TCP throughput



A simple model for TCP throughput



Implications on High-speed TCP

$$\text{Throughput} = \sqrt{\frac{3}{2}} \frac{1}{RTT \sqrt{p}} MSS$$

- **Assume RTT = 100ms, MSS=1500bytes, BW=100Gbps**
- **What value of p is required to reach 100Gbps throughput?**
 - $\sim 2 \times 10^{-12}$
- **How long between drops?**
 - ~ 16.6 hours
- **How much data has been sent in this time?**
 - ~ 6 petabits

Topics

- Basics (lectures 1–3)
- Application layer (lectures 4, 5)
 - HTTP, DNS, and CDN
- Transport layer (lectures 6–9)
 - UDP vs. TCP
 - TCP details: reliability and flow control
 - TCP congestion control: general concepts only
- **Network layer (lectures 10, 11)**
 - Data plane

Forwarding vs. routing

- **Forwarding: “data plane”**
 - Directing one data packet
 - Each router using local routing state
- **Routing: “control plane”**
 - Computing the forwarding tables that guide packets
 - Jointly computed by routers using a distributed algorithm
- **Very different timescales!**

Designing the IP header

- **Think of the IP header as an interface**
 - Between the source and destination end-systems
 - Between the source and network (routers)
- **Designing an interface**
 - What task(s) are we trying to accomplish?
 - What information is needed to do it?
- **Header reflects information needed for basic tasks**

What information do we need?

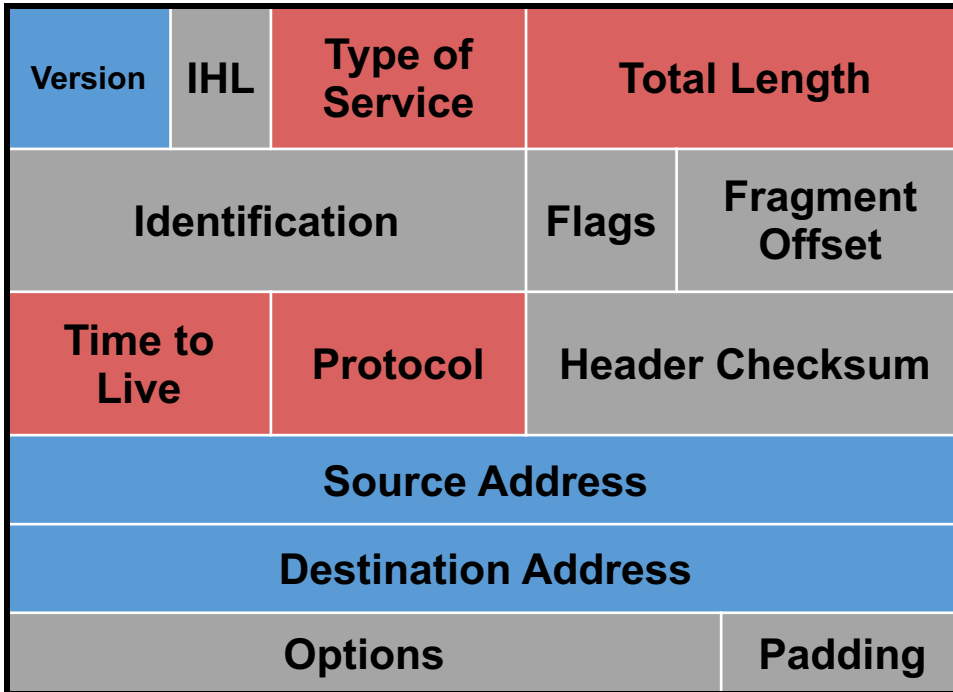
- **Parse packet**
 - IP version number (4 bits), packet length (16 bits)
- **Carry packet to the destination**
 - Destination's IP address (32 bits)
- **Deal with problems along the way**
 - Loops: TTL (8 bits)
 - Corruption: checksum (16 bits)
 - Packet too large: fragmentation fields (32 bits)

IP packet structure

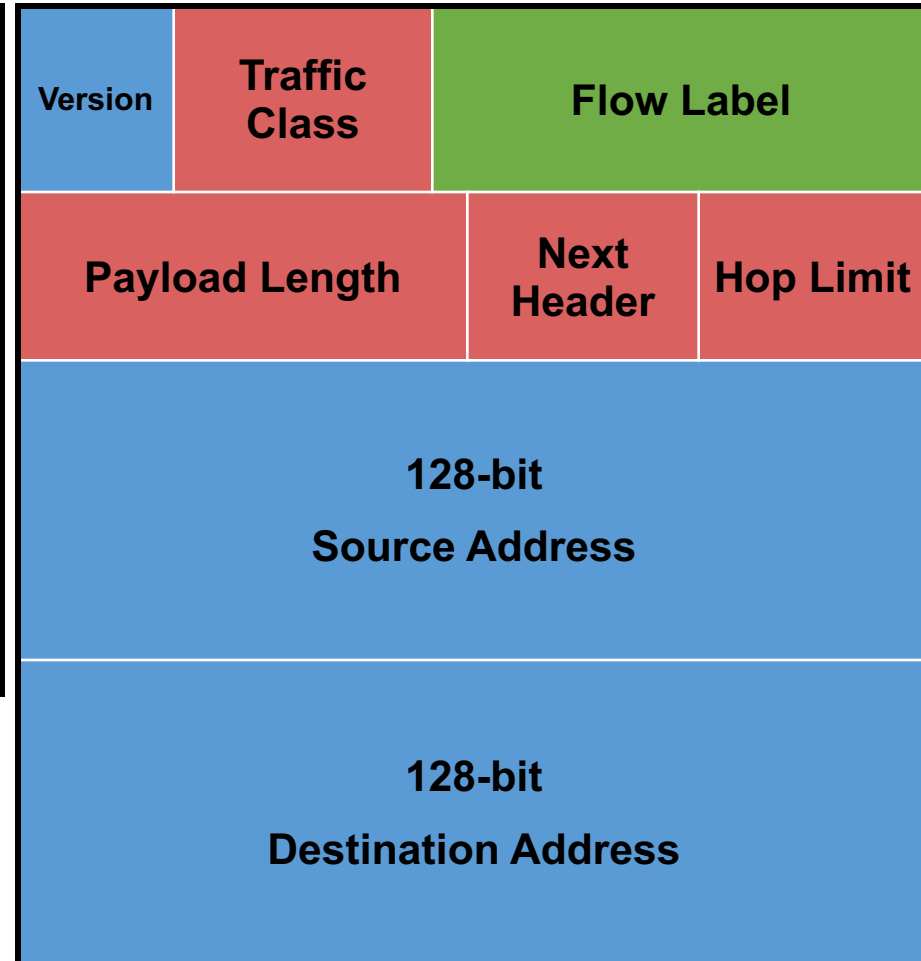
4-bit Version	4-bit Header Len	8-bit ToS	16-bit Total Length (Bytes)	
16-bit Identification			3-bit Flags	13-bit Fragment Offset
8-bit TTL		8-bit Protocol	16-bit Header Checksum	
32-bit Source IP Address				
32-bit Destination IP Address				
Options (if any)				

IPv4 and IPv6 header comparison

IPv4



IPv6

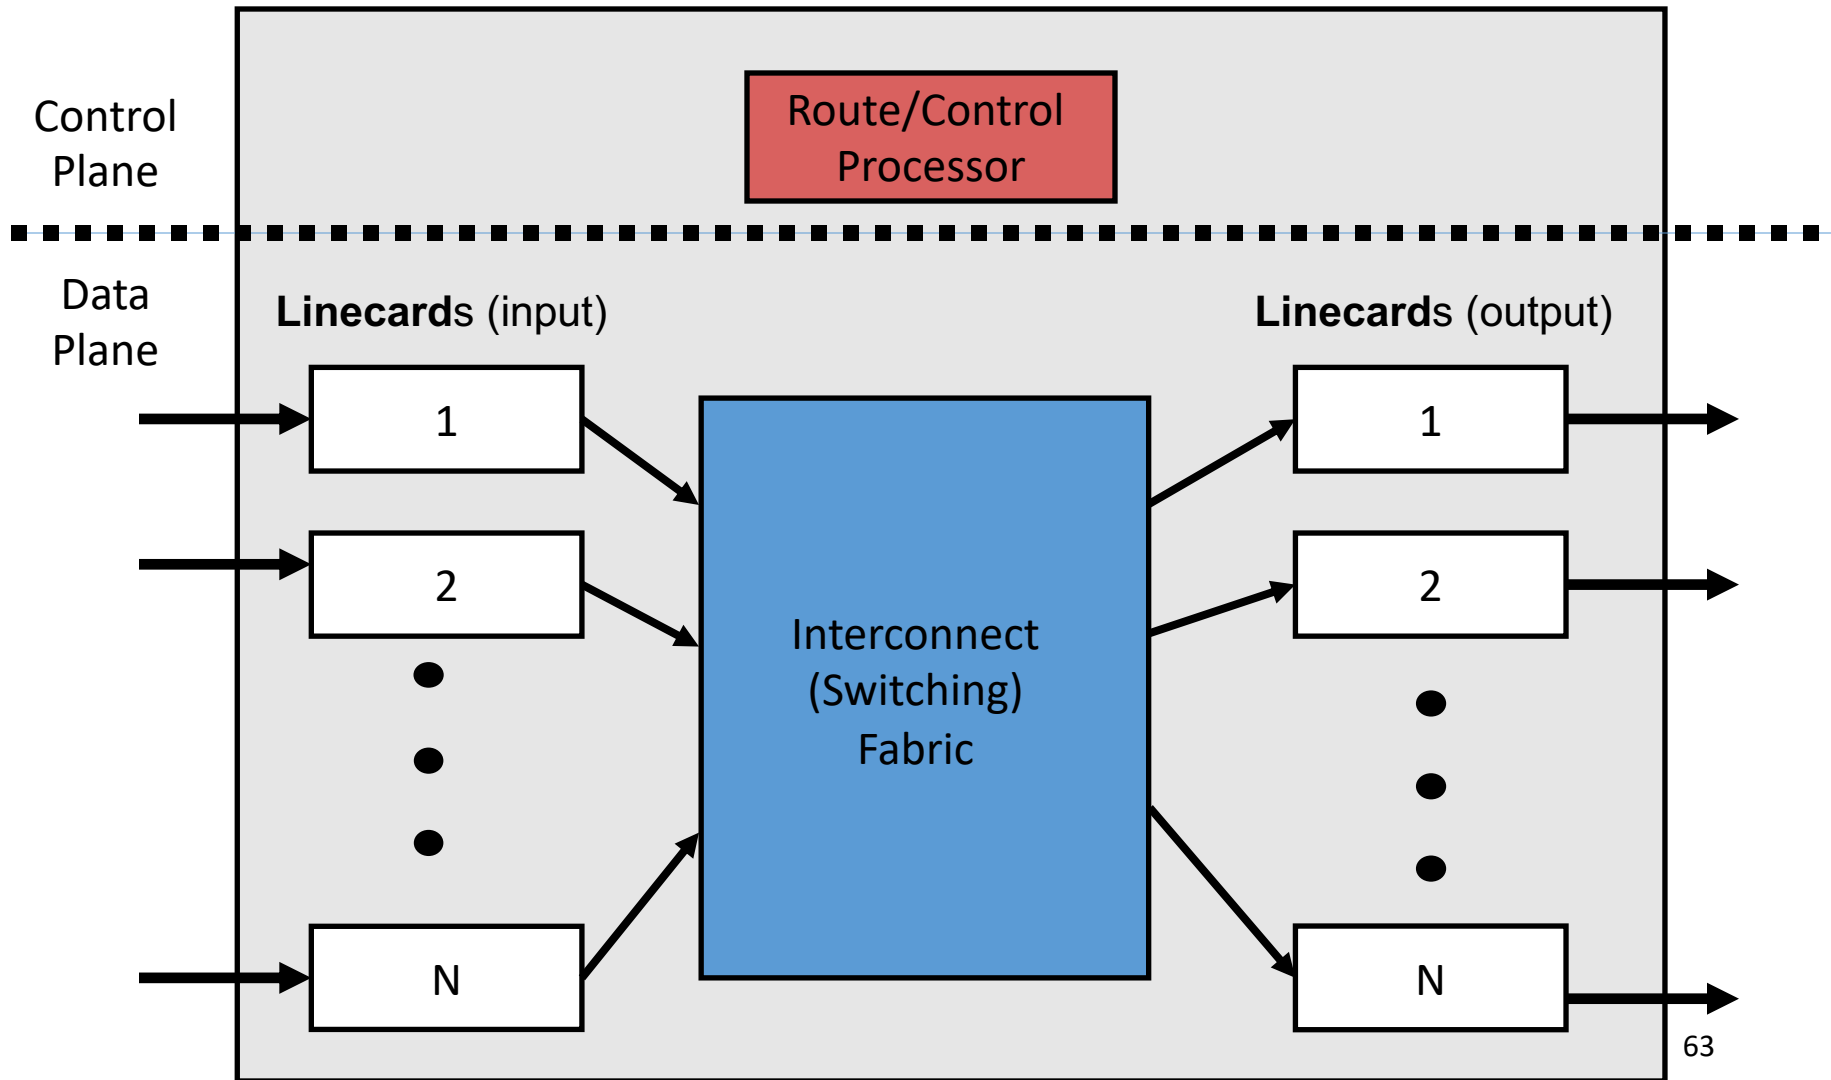


- Field name kept from IPv4 to IPv6
- Fields not kept in IPv6
- Name & position changed in IPv6
- New field in IPv6

Philosophy of changes

- **Don't deal with problems: leave to ends**
 - Eliminated fragmentation and checksum
 - Why retain TTL?
- **Simplify handling:**
 - New options mechanism (uses next header)
 - Eliminated header length
 - Why couldn't IPv4 do this?
- **Provide general flow label for packet**
 - Not tied to semantics
 - Provides great flexibility


What's inside a router?



Looking up the output port

- **One entry for each address → 4 billion entries!**
- **For scalability, addresses are aggregated**

Longest prefix matching



Send to the port with the longest prefix match

Thanks!
Q&A